



FOM Hochschule für Oekonomie & Management

Hochschulzentrum Düsseldorf

Scientific Paper

part-time degree program

5th Semester

in the study course "Wirtschaftsinformatik"

as part of the course

Big Data & Data Science

on the subject

**Predicting Music Genres based on Spotify Song Data using a Gradient
Boosting Algorithm**

by

Thomas Keiser

Martin Krüger

Luis Pflamminger

Jesper Wesemann

Advisor: Prof. Dr. Adem Alparslan

Matriculation Number: 534320 (Keiser), 534306 (Krüger), 538276 (Pflamminger), 533882 (Wesemann),

Submission: January 31st, 2022

Contents

List of Figures	IV
List of Tables	V
List of Abbreviations	VI
1 Introduction	1
1.1 Problem Definition and Goal	1
1.2 Structure and Methodology	2
2 Fundamentals	3
2.1 Classification in the Context of Big Data	3
2.2 Decision Trees	4
2.2.1 Decision Tree Algorithm	5
2.2.2 Evaluation of Decision Trees	10
2.3 Gradient Boosting	11
2.3.1 Gradient Boosting Algorithm	12
2.3.2 Evaluation of Gradient Boosting	19
2.4 Concepts for Data Preparation and Modeling	20
2.4.1 Mean Removal, Variance Scaling and Standardization	20
2.4.2 Dimension Reduction	20
2.4.3 Hyperparameter Optimization	21
2.4.4 K-fold Cross Validation	21
2.5 Cross Industry Standard Process for Data Mining	22
2.6 Web Application Programming Interfaces	23
2.6.1 Use Cases and Types of Web APIs	23
2.6.2 The HTTP Protocol	24
2.6.3 JSON	26
2.7 Basic Concepts of Music Theory	27
2.7.1 Melodical Components of Music	27
2.7.2 Lyrics and Instruments	28
2.7.3 Chronological Classification	29
2.7.4 Genres	29
3 Implementation	33
3.1 Data Collection	33
3.1.1 Requirements for the Dataset	33

3.1.2	Ressources and Approach	34
3.1.3	Getting Features	34
3.1.4	Getting Track IDs and Labels	36
3.1.5	Python Implementation	38
3.2	Data Understanding	39
3.2.1	The Streaming Service Spotify	40
3.2.2	Feature Analysis	41
3.2.3	Overall Correlation	43
3.2.4	Comparison of Music Theory to Audio Features	44
3.3	Data Preparation	48
3.3.1	Import and Cleaning	48
3.3.2	Data Transformation	50
3.3.3	Principal Component Analysis	51
3.4	Modeling	52
3.4.1	The Classifier	52
3.4.2	Validation Curves	52
3.4.3	Hyperparameter Tuning and Fitting a Model	54
3.5	Evaluation	56
3.5.1	Gradient Boosting Evaluation	56
3.5.2	Process Evaluation	60
4	Conclusion	61
	Appendix	63
	Bibliography	64

List of Figures

Figure 2: Decision tree first split	8
Figure 3: Decision Tree second Split	10
Figure 8: Description of different parts of a URL	24
Figure 9: Example HTTP request and response messages	25
Figure 10: Audio Feature Request	35
Figure 11: Artist Request	36
Figure 12: Categories and Playlists in Spotify App	37
Figure 14: Correlations inside categories	44
Figure 15: <i>loudness</i> compared in the different categories	45
Figure 16: <i>instrumentalness</i> compared in the different categories	46
Figure 17: <i>energy</i> compared in the different categories	46
Figure 18: Different features visualized in distribution- and boxplots	48
Figure 19: Test accuracy for each dimensionality after PCA	51
Figure 20: Validation Curves modulatng 3 hyperparameters	54
Figure 21: Comparison of ROC and AUC	59
Figure 22: Comparison of feature importance	60

List of Tables

Table 1: Input dataset	6
Table 2: Input data of second split	9
Table 3: Input dataset	13
Table 4: Pseudo Residuals for first Iteration	15
Table 5: Output values for first iteration	16
Table 6: Gradient boosting complete	17
Table 7: Gradient boosting Final Output	19
Table 8: Common status codes	26
Table 10: Confusion matrix	57
Table 11: Gradient boosting confusion matrix	57
Table 12: Gradient boosting classification report	58

List of Abbreviations

API	Application Programming Interface
URL	Uniform Resource Locator
JSON	JavaScript Object Notation
HTTP	HyperText Transfer Protocol
PCA	Principle Component Analysis
CART	Classification and Regression Tree
GBM	Gradient Boosting Machine
PR	Pseudo Residual
ROC	Receiver Operating Characteristic
AUC	Area Under the Curve
CRISP DM	Cross Industry Standard Process for Data Mining
CV	Cross Validation

1 Introduction

Spotify provides the world's leading streaming service. Since its launch in 2006, the company has gained over 365 million users, of which nearly 45 percent are subscribed to the chargeable premium service. The ability to listen to almost any song a user might want is a great benefit streaming services have over regular music vendors like iTunes. On the other hand, users might quickly get lost or feel overwhelmed by such a large collection to choose from. To guide users and help them find the music they want to listen to in a certain situation, Spotify uses a number of methods like premade playlists or categories for specific moods and genres. The biggest benefit to the user experience might certainly be the personalized playlists, radios or artist and song recommendations. To be able to have such a robust recommendation system, Spotify needs to understand each user's listening behaviour and have methods to predict, which music a user might enjoy depending on their past usage. Not just user data is important to gain this knowledge, but Spotify also needs to understand how to categorize music itself and find ways to identify which songs are alike and how they relate to each other. Data Science is essential to achieve this goal. Machine learning is used to analyze the music in their catalogue and create characteristics about it. Spotify offers this data to the public, which forms the foundation of the dataset.

1.1 Problem Definition and Goal

The paper is set in the context of Big Data Analytics and aims to explain and implement many of its concepts and processes. The basic goal of the project is to develop a machine learning model. During this process, general data mining stages, as layed out in the Cross Industry Standard Process for Data Mining (CRISP DM), must be theoretically explored, understood and practically implemented. This includes steps like data preparation, model creation and evaluation. Creating a model with a high accuracy is only a secondary concern in this project. Instead gaining a solid understanding of data mining concepts is paramount.

To achieve this goal Spotify song data is collected and an attempt is made to build a Gradient Boosting Algorithm, which is able to classify songs into different genres to a reasonable accuracy based on preprocessed audio features from Spotify.

1.2 Structure and Methodology

The paper is divided into four main sections, starting with the introduction which contains the problem statement, goal and overall structure.

The second section discusses fundamental concepts. This includes embedding this project into the larger context of Big Data, explaining the algorithms used on a theoretical level and introducing further concepts which are used for data preparation and modeling. Additionally, the CRISP DM model is explained, which is the basic structure that implementation is based on. The second section is concluded by an introduction to the use cases and functionality of Web Application Programming Interfaces (APIs) and a discussion of basic concepts in music theory.

The third section begins with the data collection process and continues on with the practical implementation steps of understanding the dataset's features and labels in the context of Spotify's analysis, preparing and analyzing the data for modeling, creating a gradient boosting classifier and finally evaluating the project results.

The paper concludes with a summary of the insights gained and embedding the project into a broader context.

Facts presented as part of the paper's fundamental sections are derived from literature research using renowned book sources and scientific papers. Additionally online articles were used to round out the research. Most visualizations in this paper are made by hand using Python libraries such as *seaborn* and *matplotlib*. For data collection and understanding Spotify's developer resources are used as a reference. Documentation from Python libraries such as scikit learn is used extensively during the preparation, modeling and evaluation parts of the project.

2 Fundamentals

In this section the project is embedded into the larger context of Big Data, algorithms are explained on a theoretical level and further concepts for data preparation and modeling are introduced. Additionally, the CRISP DM model, Web APIs and basic concepts of music theory are explained.

2.1 Classification in the Context of Big Data

The project falls under the umbrella term Big Data. Big Data is defined as a large set of information which is stored and processed to create business value. The analysis represents a substantial part of Big Data and is referred to as Big Data Analytics.

The term analytics is used to describe a systematical analysis of data of any format. Analytics is about detecting hidden patterns, clusters and meaningful features ranging from simple detection to deep analysis and predictions [1, p.2] [2, p.8f]. The creation of value is closely linked to the use cases for which in-depth knowledge is required. Generally, literature distinguishes between four subsections of analytics by grouping them according to the methodology and their objectives. Descriptive analytics is the simplest form of analytics and is often the first step for further research. It gives information about what has happened in the past and is mostly used for reporting. Diagnostic analytics is based on descriptive analytics but serves a different goal, since its objective is to give insights on why something has happened. Diagnostic analytics evaluates the impact of features, detects correlations and dependencies. Predictive analytics takes this approach one step further as it predicts what is most likely going to happen in the future based on the knowledge gained from past data. It creates models with the help of algorithms to determine the probability of outcomes. The final step of analytics is prescriptive analytics which is again based on its predecessor. Prescriptive analytics predicts outcomes and recommends actions to avoid or support them respectively [2, p.8f].

This project contains elements of descriptive, diagnostics and predictive analytics. Using the CRISP DM process model, the dataset is analyzed first. Afterwards a predictive model is generated to classify new data based on known patterns.

Big Data Analytics covers a large variety of research areas that differentiate mostly in their input and learning methodology.

Input data exists in various formats for which adequate storage solutions and machine learning algorithms are required. Conventional data is almost exclusively stored in structured data formats while modern developments increasingly rely on semi- and unstructured

data. Unstructured data has no fixed format, scheme or structure. Structured data, on the other hand, has a fixed format and fits into a predefined data models. Semi-structured data is an intermediate stage and is not based on a strict standard. It often consists of pre-defined tags like JavaScript Object Notation (JSON). This project utilizes semi-structured data during the collection-phases which is converted into a structured format for the modeling.

The fundamental learning methods are unsupervised, supervised and reinforced learning with semi-supervised learning as a subcategory. The methods differ in how the model learns and which input data is required [3, p. 4]. For supervised learning, a model is created, which maps input values, often referred to as targets, to a known label value. The model thus learns on already correct and complete data through the detection of hidden pattern and correlations. The target labels are either nominal or numerical which results in a classification or regression task [4, p. 46]. Unsupervised learning forms the opposite as its input data has no target labels. The algorithm is autonomous in deriving or recognizing common structures [5, p. 97]. Semi-supervised learning is set between both extremes and includes samples from both categories. Reinforcement learning is a fundamentally different learning approach with no training data being available at the beginning [3, p. 7]. The algorithm acquires data only by interacting with the environment. Models train by tackling challenging problems and having their decisions immediately rewarded if successful or punished if unsuccessful through feedback signals [5, p. 98]. This project uses a Gradient Boosting Algorithm, which belongs to the category of supervised learning methods. Therefore, the dataset consists of features and labels.

2.2 Decision Trees

Decision trees are one of the most widely used Supervised Machine Learning Algorithms either as standalone solutions or in combination with enhancement approaches like boosting. They allow for a very flexible construction and can be utilized for various machine learning problems such as classification and regression.

Decision trees “predict an unknown value of a target variable by learning decision rules from data features” to reconstruct the dependence between the features and the respective labels for each sample. To perform classification or regression, Decision Trees rely on recursive splitting of the dataset into multiple subgroups. As the number of iterations increases, the subgroups become more and more homogeneous [6, p.330]. The ideal result is that each subgroup is fully homogeneous and therefore only represents a single category (in case of classification). However, this is often only a theoretical best condition,

as multiple risks, such as overfitting, are associated with the increasing depth of Decision Trees.

Trees consist out of four main components. A node is a discrete decision function that takes samples as its input and splits them based on features into subgroups. The aim of each split, as previously discussed, is to create a split that results in the overall most homogeneous distribution for all subgroups [7, p.6]. Nodes can be subclassified into three kinds. The top-node, from which the classification starts, is called a root node. Nodes that are located at the very end of a Decision Tree are referred to as leaves. Leaves do not split data any further and only mark the end of a decision tree. When reached, leaves categorize or predict a final output value depending on the prediction task. Nodes in between the root node and leaves are called internal nodes. Like the root, internal nodes are responsible for the recursive splitting of the data. Branches connect nodes with each other [7, p.4]. For classical trees, information only flows from top to bottom of the tree.

In practice, various implementations for computing decision trees exist. Each follows the same principle of regressively finding perfect splits to separate the data, but uses different methods to find the ideal splitting criteria, which strongly influences the structure of the tree, its accuracy and performance. Thus, each implementation has benefits and constraints which have to be taken into account when determining the model. This project uses the Python library sklearn to implement a classification tree. Sklearn is based on the Classification and Regression Tree (CART) algorithm [8].

2.2.1 Decision Tree Algorithm

Initial Dataset:

To better visualize the procedure of the decision tree algorithm, a simplified dataset is used, on which the individual steps are explained. For this example, a classification problem is chosen. The dataset consists of actual features and labels from the project implementation phase. The features are *acousticness* and *danceability*. Both features are numerical with a value range in between 0 and 1. The classification problem is binary with *hiphop* and *jazz* representing the classes k for which the samples of the dataset are classified. Mathematically, the dataset is represented in the following form: x_i presents a set of explanatory features while y_i represents the corresponding label for one data point of the input dataset N with a total number of n samples.

Table 1: Input dataset

category	track	feature_danceability	feature_acousticness	label
hiphop	h1	0.949	0.132	1
hiphop	h2	0.743	0.234	1
hiphop	h3	0.913	0.394	1
hiphop	h4	0.810	0.504	1
hiphop	h5	0.434	0.198	1
jazz	j1	0.654	0.534	0
jazz	j2	0.593	0.312	0
jazz	j3	0.234	0.341	0

Recursive Tree Construction:

The decision tree algorithm splits the dataset recursively into subsets. At the beginning, the whole dataset is contained within the root node from which the division into subtrees starts. Each node is mathematically described as Q_m and contains a subgroup N_m of the input dataset. Nodes split their respective groups of data further or act as a final classifier in the form of a leaf. The split can be binary or multiway. CART only utilizes binary splits and divides the node Q_m into two repetitive subgroups labeled Q_m^{left} and Q_m^{right} [8].

The algorithm is greedy and thus determines an optimal division of the samples according to the impurity of both subgroups for each split. The goal for every split is to divide the group of data into two subgroups, which are more homogeneous than the origin and maximize the overall homogeneity. The division of the data is always performed using a feature, on the basis of which the samples can be divided either by a threshold for nominal values or an is-equal-to query for categorical values. The split is mathematically described as the following equation 1. It consists of the the feature j and a criteria t_m [8].

$$\theta = (j, t_m) \quad (1)$$

Splitting Criteria and Information Gain:

The optimal split is determined by selecting the feature from which the most information can be gathered. There are multiple approaches for finding this split with the most common one being the information gain. Ideally the gain should be maximised. A simplification for the information gain is presented in (2). To receive information gain, $G(Q_m, \theta)$ must be subtracted from the entropy of the parent node [9, p.613f]. The simplification is to find the minimum $G(Q_m, \theta)$ as stated in (3). Both approaches lead to the same result [8].

$$G(Q_m, \theta) = \frac{N_m^{left}}{N_m} H(Q_m^{left}, \theta) + \frac{N_m^{right}}{N_m} H(Q_m^{right}, \theta) \quad (2)$$

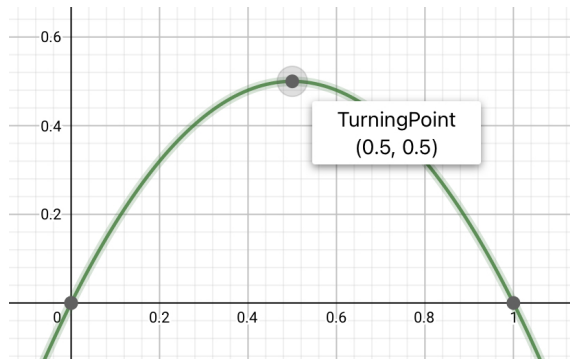
$$\theta^* = \arg \min_{\theta} G(Q_m, \theta) \quad (3)$$

To determine the overall best $G(Q_m, \theta)$ the respective impurities H for both subgroups Q_m^{left} and Q_m^{right} are required. Each impurity is weighted according to its relative size $\frac{N_m^{left}}{N_m}$.

The impurity H for CART is calculated by the gini index (4) [9, p.613f]. Gini index measures the probability that a sample does not belong to the category that represents the majority of the subgroup [6, p.335]. If both categories of a subgroup are identical in size, the gini index reaches its maximum point at 0,5 (figure 1). The maximum of the gini index means the worst possible data constellation for a subgroup with maximum heterogeneity. Gini index equal to 0, on the other hand, represents the best possible result with the subgroup being fully homogenous. p_i represents the probability that a sample belongs to the class j [6, p.335].

$$Gini = 1 - \sum_{j=1}^k (p_j)^2 \quad (4)$$

Figure 1: Gini Graph



The optimal splits for the first iteration of the dataset look like the following. For numerical features the determination of the optimal split is more complex as for categorical ones because the calculation is not straight forward but must be computed for every value in the value range to find the best overall split for a single feature. The example only shows the best splitting criteria while a trial and error approach can be found in the appendix.

Danceability:

$$Gini^{left} = 1 - \left(\left(\frac{2}{5}\right)^2 + \left(\frac{3}{5}\right)^2\right) = 0,48$$

$$Gini^{right} = 1 - \left(\left(\frac{3}{3}\right)^2 + \left(\frac{0}{3}\right)^2\right) = 0$$

$$G(Q_0, \theta) = \frac{5}{8} * 0,48 + \frac{3}{8} * 0 = 0,30$$

Acousticness:


$$Gini^{left} = 1 - \left(\left(\frac{4}{4}\right)^2 + \left(\frac{0}{4}\right)^2\right) = 0$$

$$Gini^{right} = 1 - \left(\left(\frac{1}{4}\right)^2 + \left(\frac{3}{4}\right)^2\right) = 0,36$$

$$G(Q_0, \theta) = \frac{4}{8} * 0 + \frac{4}{8} * 0,36 = 0,18$$

The conclusion is that the feature *danceability* creates a better initial prediction than *acousticness* does. It is thus determined to be the initial splitting criterion for the classification tree (figure 2). With the first iteration completed, the next iteration will start with each leaf as its input. The second iteration will split the nodes further if no stop criterion is met.

Figure 2: Decision tree first split



Stop Criteria:

The natural stop criterion is a completely homogeneous data-group for a node. If a node only contains one single class of samples, no further splitting is possible. This is the case for the right node of the example marked in blue. The samples are completely homogeneous and only consist of the category *hiphop*. The node therefore automatically becomes a leaf and represents the category of samples. Other stop criteria can be predefined depending on own preference. The most relevant criterion is the definition of a maximum depth of the decision tree. Maximum depth describes a maximum amount of recurrent

iterations before the algorithm automatically stops and treats the last nodes as leaves [7, p.7]. Maximum depth belongs to a set of hyperparameters that will be discussed in section 3.4.

Completion of the example:

The second iteration has both nodes created by the first iteration as its input but only the yellow node requires further splitting as the blue node is already fully classified. For the second iteration only the feature *acousticness* is relevant as *danceability* cannot split this subgroup of the dataset further. The input data for the second split looks like the following (table 2) and the same calculation takes place again.

Table 2: Input data of second split

category	track	feature_danceability	feature_acousticness	label
hiphop	h5	0.434	0.198	1
jazz	j1	0.654	0.534	0
jazz	j2	0.593	0.312	0
jazz	j3	0.234	0.341	0

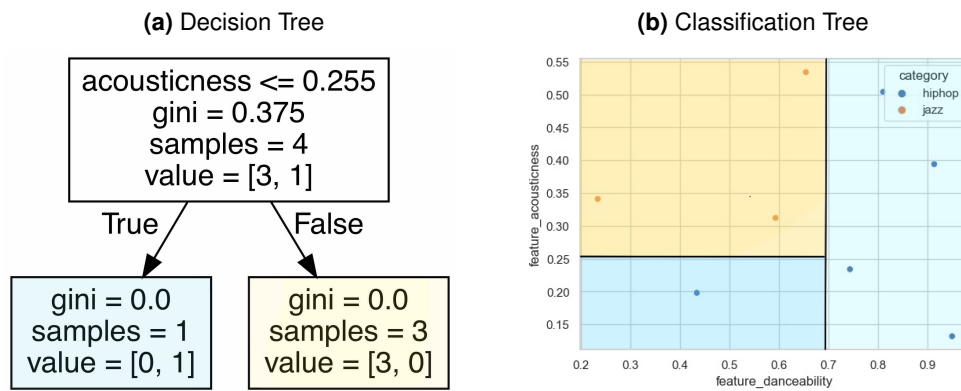

Acousticness:

$$Gini^{left} = 1 - \left(\left(\frac{0}{1}\right)^2 + \left(\frac{1}{1}\right)^2\right) = 0$$

$$Gini^{right} = 1 - \left(\left(\frac{3}{3}\right)^2 + \left(\frac{0}{3}\right)^2\right) = 0$$

$$G(Q_1, \theta) = \frac{1}{4} * 0 + \frac{3}{4} * 0 = 0$$

This time *acousticness* leads to two fully homogeneous subgroups and therefore marks the end of the classification tree as all samples are classified correctly. The second split looks like the following with the total classification tree displayed in figure 3.

Figure 3: Decision Tree second Split**Figure 4: Final Classification Tree**

2.2.2 Evaluation of Decision Trees

In conclusion, decision trees can be assessed as follows. Starting with the advantages, the main benefit is the overall simplicity of decision trees, both from a technical and business point of view [6, p.339]. For researchers and developers, trees are easy to construct,

require little to no data preparation, are almost universally applicable with the possibility of validation. However, the simplicity for business should not be underestimated either. When comparing machine learning algorithms, the main comparison is often the accuracy of a model. The areas in which decision trees stand out include visualization and comprehensibility. The decision tree algorithm is a white box model that allows complete transparency and explainability [8, p. 10.10.].

The disadvantages of decision trees are again closely related to their simplicity. Overfitting and the relative instability of decision trees are the main drawbacks and result in good memorization but a comparatively weak generalization ability [6, p.340] [8, p. 10.10.].

2.3 Gradient Boosting

The Gradient Boosting Algorithm is derived from Gradient Boosting Machine (GBM), which are a family of powerful machine learning algorithms with a certain procedure pattern for the creation of models. In general, GBMs are very flexible in their characteristics with the possibility of utilizing multiple different Machine Learning algorithms as their foundation [10].

Boosting differs from classical approaches as it does not consist out of a single predictive model but an ensemble approach. Ensemble algorithms contain multiple weak learners that form a committee to create a strong prediction. Weak learners are often very simple forms of traditional algorithms, like decision trees, and must just be able to predict parts of the dataset correctly. Only the combination of many weak learners allows the model to perform overall accurate predictions [11] [12, p.1f]. The most common form of ensemble algorithms are bagging algorithms with random forests as an example. Bagging, in essence, is the combination of multiple unique models. The prediction is formed by aggregating the outputs from all models into a single representative value. Typically, all models are derived from a single algorithm, like decision trees for random forests, but technically there is no limitation to aggregate outputs from different algorithms. Which is also true for other ensemble algorithms [12, p.2].

Boosting, on the other hand, follows a different principle and does not rely on independent models with an aggregation function. Boosting fits new models sequentially and can thereby use earlier acquired knowledge for further iterations. This allows a GBM to train specific areas of the dataset where it has previously performed poorly [12, p.11] [6, p.345f].

2.3.1 Gradient Boosting Algorithm

The generic Gradient Boosting Algorithm follows a sequence of three steps beginning with the initialization of the dataset, to the sequential training and ending with a final output. At the beginning, an additional initiation of the dataset and a loss function is necessary [13]. The mathematical representation of the dataset is like the one used for decision trees. A summary: x_i represent the explanatory features while y_i represents the corresponding label for one data point of the input dataset N with a total number of n samples.

The mathematical goal of the algorithm is to reconstruct the unknown functional dependence f between x_i and y_i with an estimate F^* for every data point, such that the specific loss function $L(y, \gamma)$ is minimized (5) [14, p.1189] [10, p. 2.1].

$$F^* = \arg \min_{\gamma} L(y, \gamma) \quad (5)$$

Loss Function and Weak Learners:

The loss function is an indicator for the quality of the model. A small loss for a data point means that the prediction is close or identical to the observed label and the model therefore categorizes the sample correctly whereas a high loss implies that the model could not predict the sample well. Given a particular learning task and dataset, different loss functions must be considered as loss functions are only suitable for specific data and task constellations. The most common loss function for binary classification is the so-called bernoulli loss (6). The bernoulli loss can be transformed into a $\log(odds)$ -prediction (7) as it is better suited for further calculations [10, p. 3.1] [15]. Variations of (7) will be used in the following section to demonstrate the gradient boosting procedure.

$$L(y_m, \gamma) = -\log(\text{likelihood}) \quad (6)$$

$$L(y_m, \gamma) = -y_i * \log(odds) - \log(1 - p) \quad (7)$$

Additionally, a machine learning algorithm must be defined as a weak learner. For GBMs there are multiple learners to choose from. Again the choice mostly depends on the prediction task and available data [10, p. 3.2]. A classical approach is the use of Decision Trees, which was also chosen as the weak learner for this project. Decision Trees used for Gradient Boosting are always Regression Trees, regardless of whether they are used for regression or classification problems. The optimization parameters are almost identical

to the ones of standalone Decision Trees, but the trees often look very different because they are specifically created as weak learners. As a result, the Decision Trees often only consist of very few layers with only 8-32 leaves.

Initialization:

To showcase the Gradient Boosting Algorithm the same sample dataset is used as for Decision Trees. It again consists of eight samples with two features and two categories as target labels.

Table 3: Input dataset


category	track	feature_danceability	feature_acousticness	label
hiphop	h1	0.949	0.132	1
hiphop	h2	0.743	0.234	1
hiphop	h3	0.913	0.394	1
hiphop	h4	0.810	0.504	1
hiphop	h5	0.434	0.198	1
jazz	j1	0.654	0.534	0
jazz	j2	0.593	0.312	0
jazz	j3	0.234	0.341	0

The first step is to set an initial prediction for all samples of the dataset. The initial prediction is not unique for individual samples but a uniform value. The optimal initial prediction can be calculated using the following equation (8) [16, p.361]. For $F_0(x)$, representing the initial prediction, a minimum of γ is searched for. The right-hand side of the equation only consists of a sum for each sample i (of the total dataset N) of the known loss function with the respective label y_i and γ as its input. To find the low point of the equation, the derivate of the loss function is required. The final calculation of the overall $\log(odds)$ that a song is classified as *hiphop* is the \log_e of the sum of songs of the category *hiphop* divided by the sum of the songs of the category *jazz*. The result can be checked graphically as is equal to the x_1 -coordinate value of the low point of the $\log(odds)$ -prediction (figure 5).

$$F_0(x) = \arg \min_{\gamma} \sum_{i=1}^n L(y_i, \gamma) \quad (8)$$

$$F_0(x) = \log_e\left(\frac{5}{3}\right) = 0.51$$

Figure 5: Gradient boosting minimum for first prediction



Sequential Processing:

With the completion of the initialization of the dataset the modeling can begin. The training is a sequential process of constructing regression trees with a total of M iterations. The first iteration starts with $m = 1$. The modelling consists out of four sub-steps which are numbered consecutively [14, p.1198].

First, the $\log(odds)$ -prediction must be converted back into a probability p with the help of a logistic function as probability is easier to use for classification (9). The result is that all songs have a 63% chance of belonging to the category *hiphop*.

$$p = \frac{e^{\log_e(odds)}}{1 + e^{\log_e(odds)}} \quad (9)$$

$$p = \frac{e^{\log_e(\frac{5}{3})}}{1 + e^{\log_e(\frac{5}{3})}} = 0,63$$

In the **first step** the Pseudo Residuals (PRs) r_{im} for each sample i of the dataset are created. The equation (10) for calculating the PR consists out of known fragments. For every sample i a PR r_{im} is calculated using the derivative of the $\log(odds)$ -prediction with the label y_i and the prediction of the last iteration $F = F_{m-1}$ as its input [16, p.361]. Again, the equation can be simplified greatly. For each sample the PR can be calculated by only subtracting the previously calculated probability p from the observed label y [13] [17]. Ideally, an additional column is created in which the PRs are temporarily stored (table 4).

$$r_{im} = -\left(\frac{\partial L(y_i, F(x_i))}{\partial F(x_i)}\right)_{F=F_{m-1}} \quad (10)$$

$$r_{im} = (y_i - p_i)$$

This and every following calculation is explained for the track *h1* of the dataset while the result for every other sample will be shown in form of a table.


$$r_{1,1} = 1 - 0.625 = 0.375$$

Table 4: Pseudo Residuals for first iteration

category	track	label	$F_0(x)$	probability $F_0(x)$	pseudo residuals
hiphop	h1	1	0.510826	0.624988	0.375012
hiphop	h2	1	0.510826	0.624988	0.375012
hiphop	h3	1	0.510826	0.624988	0.375012
hiphop	h4	1	0.510826	0.624988	0.375012
hiphop	h5	1	0.510826	0.624988	0.375012
jazz	j1	0	0.510826	0.624988	-0.624988
jazz	j2	0	0.510826	0.624988	-0.624988
jazz	j3	0	0.510826	0.624988	-0.624988

The **second step** constructs the a Regression Tree out of the features from the samples with the corresponding PR as the label. For this example, only tree stumps are created to simplify the implementation. The regression tree for the first iteration is shown in figure 6. After the completion of the tree, terminal regions R_{jm} must be defined for every leaf. j starts with 1 and is increased for every leaf [14, p.1195].

Figure 6: Regression tree for first iteration



In **step three**, following the completion of the regression tree, output values $\gamma_{j,m}$ are calculated by using the equation presented in (11). For each leaf in the tree, $\gamma_{j,m}$ is computed by finding $\gamma_{j,m}$ that minimizes the loss function (11) [16, p.361]. Like in the initialization step, the derivative has to be created and set equal to 0. Again, after a complicated transformation, a very simple equation remains (11). The $\gamma_{j,m}$ can be calculated using only the PR and the most recently predicted probabilities p for all samples in the leaf.

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(y_i, F_{m-1}(x_i) + \gamma) \quad (11)$$

$$\gamma_{jm} = \frac{\sum r_{im}}{\sum p_i * (1 - p_i)}$$

For $R_{2,1}$, which contains track $h1$, the calculation looks like the following. Output values of the complete dataset are displayed in table 5.

$$\gamma_{2,1} = \frac{0.375 * 4}{(0.625 * (1 - 0.625)) * 4} = 1.6$$

Table 5: Output values for first iteration

category	track	label	$F_0(x)$	pseudo residuals	output value
hiphop	h1	1	0.510826	0.375012	1.600032
hiphop	h2	1	0.510826	0.375012	1.600032
hiphop	h3	1	0.510826	0.375012	1.600032
hiphop	h4	1	0.510826	0.375012	1.600032
hiphop	h5	1	0.510826	0.375012	-1.599926
jazz	j1	0	0.510826	-0.624988	-1.599926
jazz	j2	0	0.510826	-0.624988	-1.599926
jazz	j3	0	0.510826	-0.624988	-1.599926

Step four marks the end of the first iteration and creates a new prediction $F_m(x)$ for each sample. The new $\log(odds)$ -prediction is based on the last $\log(odds)$ -prediction plus the learning rate ν multiplied by the output values for the sample of the last regression tree (12) [14, p.1203]. Normally, there is only one output value for a sample which makes the summation sign obsolete. The learning rate ν is a hyperparameter for gradient boosting. For this example, a high ν of 0.6 is used to better visualize the changes. In practice a ν in the order of 0.1 is common as a small learning rate tends to give better results but also requires much more iterations [14, p.1206].

$$F_m(x) = F_m(x-1) + \nu * \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm}) \quad (12)$$

The prediction $F_1(x)$ for $h1$ is shown below while the complete dataset is again displayed in table 6 with also the probabilities shown. As one can see, the probabilities for most samples got closer to the label. However, for track $h5$ the prediction got worse. This is why multiple iterations are necessary to form a well-founded prediction.

$$F_1(x) = 0.51 + 0.6 * 1.6 = 1.47$$

Table 6: Gradient boosting complete

category	track	label	$F_0(x)$	probability $F_0(x)$	$F_1(x)$	probability $F_1(x)$
hiphop	h1	1	0.510826	0.624988	1.470845	0.813163
hiphop	h2	1	0.510826	0.624988	1.470845	0.813163
hiphop	h3	1	0.510826	0.624988	1.470845	0.813163
hiphop	h4	1	0.510826	0.624988	1.470845	0.813163
hiphop	h5	1	0.510826	0.624988	-0.449130	0.389579
jazz	j1	0	0.510826	0.624988	-0.449130	0.389579
jazz	j2	0	0.510826	0.624988	-0.449130	0.389579
jazz	j3	0	0.510826	0.624988	-0.449130	0.389579

The second iteration of the modelling phase follows the exact same principles and will again be showcased for track *h1*. It marks the end for this example.

- New probability:

$$p = \frac{e^{1.47}}{1 + e^{1.47}} = 0,81$$

- New pseudo residuals:

$$r_{1,2} = 1 - 0.81 = 0.19$$

- New Decision Tree with terminal regions (figure 7)
- New output value:

$$\gamma_{2,2} = \frac{0.98}{0.54} = 1.81$$

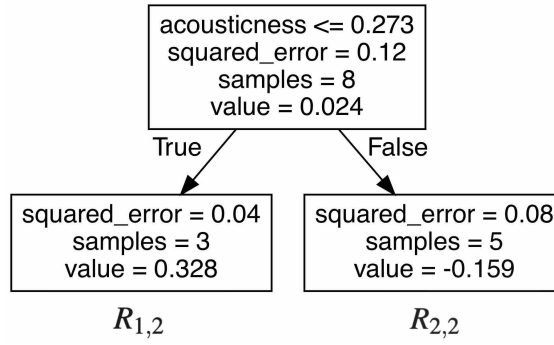
- New prediction:

$$F_2(x) = 1.47 + 0.6 * 1.81 = 2.56$$

- New probability:

$$p = \frac{e^{2.56}}{1 + e^{2.56}} = 0.93$$

Figure 7: Regression Tree for second iteration



The iterative modeling is repeated until M is reached. M marks the completion of training for the Gradient Boosting model. $F_M(x)$ is the final prediction for every sample. Lastly, the predictions must again be transformed to probabilities like for the calculation of the PR. For the final probabilities, thresholds are used to compute the category to which the samples belong [14, p.1204]. A typical threshold for binary classification is 0.5 as it splits $F_M(x)$ into two equal classes.

Output:

Sample $h1$ has a final probability 0.93 of belonging to the category *hiphop*. 0.93 is by far greater than the predefined threshold of 0.5 and thus $h1$ is classified as a track that belongs to the category *hiphop*. Table 7 shows that all samples are classified correctly but some less clear than others. A real-world implementation would consist out of more iterative models with different regression trees and a different learning rate to enhance the classification potential. Also the dataset would be very different with more samples that consist out of more features.

Table 7: Gradient boosting Final Output

category	track	prob. $F_0(x)$	$F_1(x)$	prob. $F_1(x)$	$F_2(x)$	Output Probability
hiphop	h1	0.624988	1.470845	0.813163	2.560923	0.928286
hiphop	h2	0.624988	1.470845	0.813163	2.560923	0.928286
hiphop	h3	0.624988	1.470845	0.813163	1.001911	0.731414
hiphop	h4	0.624988	1.470845	0.813163	1.001911	0.731414
hiphop	h5	0.624988	-0.449130	0.389579	0.640948	0.654953
jazz	j1	0.624988	-0.449130	0.389579	-0.918064	0.285372
jazz	j2	0.624988	-0.449130	0.389579	-0.918064	0.285372
jazz	j3	0.624988	-0.449130	0.389579	-0.918064	0.285372

An unknown sample gets initialized by $F_0(x)$ and is sequentially routed through every model. For each iteration the prediction is updated using the output value of the regression tree. Finally, the last probability is used to classify the sample using the predefined threshold.

2.3.2 Evaluation of Gradient Boosting

Gradient Boosting is a very powerful method as it can effectively capture complex dependencies for various machine learning problems. GBMs improve many existing algorithms, such as decision trees, because many problems associated with single large models are (partially) solved by the iterative ensemble approach.

The main benefit over Decision Trees is the stability of Gradient Boosting. While large trees always have to make tradeoffs between detail and overcategorization, Gradient Boosting can successively get to deeper levels of detail thanks to small trees with overall better generalization. Furthermore the flexibility of Gradient Boosting and boosting in general is massive as it only represents the framework with many parameters to adapt the algorithm very specifically to the usecase [10, p. 7.2]

The drawbacks of Gradient Boosting often arise in practice. Gradient Boosting has a significantly higher memory consumption and build time, because the model must be constructed sequentially. Also the evaluation is more time consuming as the sample must be processed by each model. From a business perspective Gradient Boosting also has disadvantages. While the overall prediction is better, it is more complex to evaluate and explain [10, p. 7.2] [12, p.27].

2.4 Concepts for Data Preparation and Modeling

In this section, concepts and processes, which are applied during data preparation and modeling, are discussed theoretically. The practical implementation is shown in sections 3.3 and 3.4.

2.4.1 Mean Removal, Variance Scaling and Standardization

During the data preparation stage of model development, it can be helpful to transform the data into a different shape to improve the results of a model [18, p. 35]. This subsection explains some methods of data transformation, which are used in the implementation section.

Mean Removal is the process of shifting the data in a column, so that the mean \bar{x} of all datapoints in the column is equal to 0 [19]. An example set of integers $a_1 = \{6, 7, 2, 1\}$ has a mean of $\bar{x}_1 = 4$. Removing its mean without distorting other properties of the data entails subtracting 4 from each member of the set, resulting in a new set $a_2 = \{2, 3, -2, -3\}$ with $\bar{x}_2 = 0$. The dataset is now centered around 0.

Variance Scaling involves transforming the data in a way, that each column has unit variance, meaning variance $\sigma^2 = 1$ [19]. This is done by dividing each member of the column by the column's standard deviation. An example set of integers $b_1 = \{6, 7, 2, 1\}$ has standard deviation $\sigma_1 = 2.94$ and variance $\sigma_1^2 = 8.67$. Dividing each member of b_1 by σ_1 results in the set $b_2 = \{2.04, 2.38, 0.68, 0.34\}$ with variance $\sigma_2^2 = 1$.

Standardization involves performing both mean removal and variance scaling on a column [19]. Standardizing an example set $c_1 = \{6, 7, 2, 1\}$ would result in the following set $c_2 = \{0.68, 1.02, -0.68, 0.102\}$. Standardized features are normally distributed. Many Machine Learning Algorithms have improved results when receiving standardized input data. It has the additional benefit of equalizing the impact of features while keeping valuable information about the outliers and value range [18, p. 35].

2.4.2 Dimension Reduction

Dimension reduction refers to reducing the number of dimensions in the input data, while keeping the highest possible amount of information from the original dataset. The benefit of this is that some features might not contribute to a better model, but instead promote bias and overfitting. There are multiple ways to reduce dimensions, such as feature selection,

where some features, which don't contribute to a better model, are simply removed. New features can be constructed by using data from different original features and combining it into one, also reducing dimensionality. Another approach is Principle Component Analysis (PCA). Here, features are constructed by finding principal components in a dataset. This is done by projecting the data into m-dimensional space, where each feature is contained in one spatial dimension. The first principal component is found by rotating the feature space to find the direction which offers the highest variance in the whole dataset. The second component must be orthogonal to the first one and again have the highest variance of all possible directions. This can be repeated to create as many principal components as are beneficial to the model. Dimension reduction using PCA was attempted as part of this project, but is not explained deeper as it didn't yield great result for the dataset.

2.4.3 Hyperparameter Optimization

Machine learning models have two types of parameters: model parameters and hyperparameters. Model parameters are values, which are optimized using training data during the training process. They are not specified in advance. Hyperparameters, on the other hand, are not optimized using data, but are specified before training. They are like "settings" for the estimator.

Even though hyperparameters are not optimized during training, but still influence the results of a model, there are multiple ways to find the best set of hyperparameters. In this project, grid search is used, which takes a parameter grid containing multiple values for each hyperparameter as input. A model is trained using every possible combination of all given values on the grid and a score is calculated for each model. The best combination of hyperparameters are the ones used on the model with the highest score. This might not be the optimal set of hyperparameters though, as only values in the search grid are considered.

2.4.4 K-fold Cross Validation

When training a model, it is necessary to split the data into train and test samples. Using the same data for training and calculating the accuracy score of a model would result in a model, that is fitted exactly to the training data. It will perform very well on that data, but will fail to make predictions on any data it has not yet seen. This is called overfitting.

When optimizing the hyperparameters of an estimator, one could use the test set to score the model and find the best parameters. A problem with this approach is, that now test

data, which should be completely independent from the training process, is involved in the optimization process, which could lead to the hyperparameters being "fitted" to the test data.

Cross Validation (CV) is a solution to this problem. During optimization, the training set is split into k groups, named folds. A given estimator is trained on $k - 1$ folds and its accuracy is calculated using the samples in the remaining folds. This is done k times, using a different fold for scoring each time. The average of all scores is the cross-validation score, giving a good representation of the overall performance of the estimator. The benefit of this is, that test data is kept completely separate from training data, but there is still a good way of evaluating the performance of an estimator trained with a given set of hyperparameters. This is crucial during hyperparameter optimization, as the best parameters can not be found without a good way to score each model.

2.5 Cross Industry Standard Process for Data Mining

CRISP DM is an organizational model, which provides an overview of the life cycle of data mining projects. This cycle is divided into six phases: business understanding, data understanding, data preparation, modeling, evaluation and deployment. The order of the phases is not strictly adhered to. In certain projects, the outcome of each phase determines which phase or which specific task in a phase must be performed next [20, p. 528]. However, data mining is not necessarily complete once a solution is implemented. After deployment, new insights are gained from the models results and additional data collection, which can be used to further improve the implementation [20, p. 529].

The process starts with business understanding, which is an in-depth analysis of business objectives and requirements. These insights allow an assessment of the current situation and the definition of goals [21]. With business understanding complete, data is collected during the data understanding process. The collected data could come in the form of existing database entries, answers from questionnaires or machine log entries, to name a few examples. The gathered data is aggregated and examined for usability. Insufficient or bad data can cause a model to produce unsatisfying or even misleading results [22]. In the data preparation step, the data is selected, cleaned, formatted and preprocessed to ensure a high quality dataset for modeling. This is usually the most costly and time intensive phase of the project, but is crucial for its overall success [22].

In the next stage, an appropriate model must be developed to generate a result, which satisfies the requirements defined during business understanding. During modeling, it

is important to constantly present interim results to the management using meaningful visualizations and subsequently improve the model using further input.

During the step of evaluation, the results of the final model are evaluated in the context of business intentions. This might lead to new insights for the organization and in turn alter the goal. This is an iterative process ending in a final decision to deploy a certain model [20, p. 530]. Depending on the requirements, the deployment phase could involve creating simple reports, altering systems or complex data visualizations with a constant reiteration of the data mining process [20, p. 532].

2.6 Web Application Programming Interfaces

This section gives an overview over the basic concepts and technologies behind Web APIs.

2.6.1 Use Cases and Types of Web APIs

An API in general is an interface between two pieces of software[23, S.1]. It abstracts complex functionality away and provides a simple interface to interact with to retrieve data and services[24]. Web APIs are a specific form of API that can be accessed over the internet using the HTTP protocol.[25] They are used to power desktop, web and mobile applications. pass data between different services, systems and devices. They make automated data access and retrieval possible and enable the integration of internal and external systems into processes and data flows.[26]

There are different types of Web APIs, mainly Open/Public APIs, Internal APIs and Partner APIs[25]. Open APIs are provided by a company, government or other organization and make it possible for any developer to access that organizations data or services[25]. Internal APIs are only used to share resources within one organization, while partner APIs are exposed to the public, but are only open to specific users, which might pay for the right to use the API[25].

Spotify offers a Public Web API, which is used in this project in the context of data collection (section 3.1). In order to understand the specifics of data collection, some concepts regarding the Hypertext Transfer Protocol (HTTP) protocol are explained in the following section.

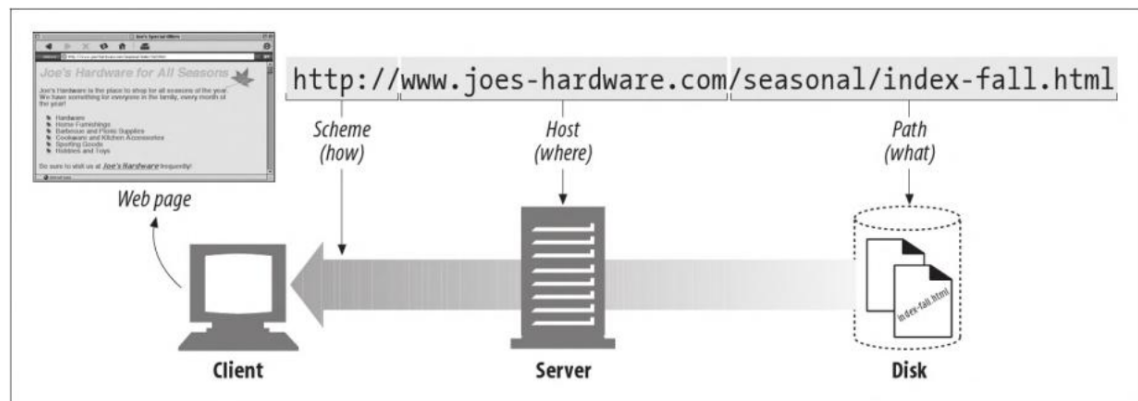
2.6.2 The HTTP Protocol

All communication on the world wide web is conducted through HTTP because of its reliability and guarantee for data integrity[27, 3f.].

Most communication on the Web is conducted between a client and a server. The client usually requests a resource from the server and the server responds with that resource, which the client can then use[27, p. 4].

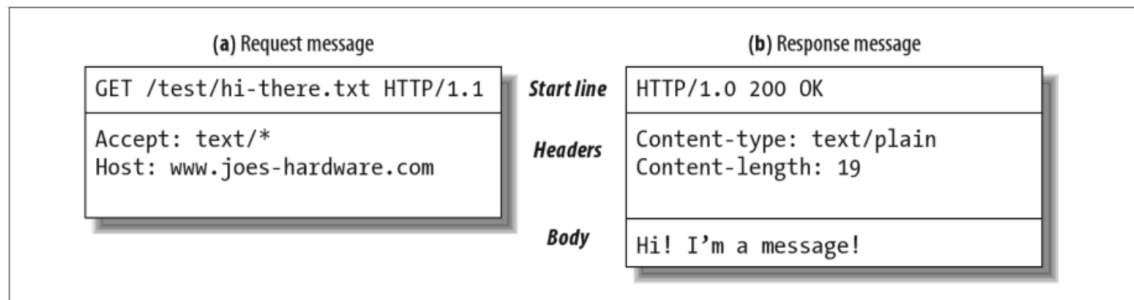
To tell the server, which specific resource it wants, the client uses a Uniform Resource Locator (URL), which uniquely identifies a resource on the web. Figure 8 shows how URLs are structured. If the server knows where to find the resource and the client is authorized to see it, the server will send it to the client. In the context of APIs, a URL pointing to a specific resource is also called an "API endpoint".[28] This interaction between client and server is usually done in an HTTP transaction.

Figure 8: Description of different parts of a URL



Source: [27, p. 24]

A transaction consists of a request (sent from client to server) and a response in which the server answers the clients request in some way.[27, p. 8] In the context of Web APIs, a request made to an API is also called "API call".[25] Figure 9 shows the structure of an HTTP request and response message.

Figure 9: Example HTTP request and response messages

Source: [27, p. 47]

Every HTTP request message starts with a method. There are many HTTP methods defined, that can be used to delete a resource, store data on the server, get header information for a document and more.[27, p. 48] In this paper, only two HTTP methods are used, GET and POST. A GET request is used to ask the server to send a resource to the client. It does not contain a body.[27, p. 48] A POST request message is used to send data to the server for processing. The data to be processed is sent in the request body.[27, p. 48]

A request also includes the URL, protocol information and a set of headers, delivering additional information about the client [27, p. 47].

Response messages contain a status code instead of a method or URL.

Every HTTP request message has a start line. It contains a method, which tells the server what action to perform on the resource, the URL and some protocol information.[27, p. 47] After the start line, some headers are given. These contain information about the request, e.g. what types of data the client can accept or credentials for authorization. Some requests also contain a body to give some data to the server. [27, p. 52]

A response message does not contain a method or URL. Instead it contains protocol information and status information in the form of a numerical status code and status text.[27, p. 48] A response also contains headers and might contain a body.[27, p. 52]

Status codes tell the client if the transaction was successful or what type of problem occurred[27, p. 49]. There are many status codes defined in the HTTP standard, some of the most common ones are shown in table 8.

Table 8: Common status codes

Status code	Reason phrase	Meaning
200	OK	Successful transaction. Any requested data is in the response body
401	Unauthorized	The response needs to contain credentials to access the requested resource.
404	Not Found	The server cannot find a resource for the requested URL.

Source: [27, p. 50]

2.6.3 JSON

The Spotify API uses the JSON data format to transfer data to the client[29].

Most modern APIs transmit data to the client using JSON. It is a text-based format able to represent semi-structured data. [30] JSON objects are simple, have good readability and are easily processed by computers, which explains the wide usage of the data format[31]. JSON is tightly integrated in the JavaScript programming language and can be used without any parsing or serialization.[31] This is beneficial, as JavaScript is widely used both on the client and server side of web applications. JSON can also be used with many other programming languages. [32]

JSON supports only six very basic datatypes: String, Number, Boolean, Null, Object and Array.[31] The following code snippet shows how each of the datatypes are defined in JSON.

```

1  {
2      "string": "Martin",
3      "number": 200,
4      "number2": 300,
5      "boolean": true,
6      "booleanFalse": false,
7      "nullValue": null,
8      "object": {
9          "objectAttribute": "string",
10         "objectAttribute2": "string2"
11     },
12     "array": [
13         "entry1",
14         "entry2",
15         {
16             "objectInArray": true
17         }
18     ]
19 }
```


The code snippet shows, that the whole JSON object is wrapped in curly braces. Each attribute is defined as an attribute-value-pair, where the attribute name is defined in double quotes, followed by a colon and the value definition. String values are defined in double quotes, numbers and booleans without. Objects are defined in curly braces and can again contain any datatype. Arrays are defined using brackets. They can contain any number of entries of any type. Types can also be mixed within the same array. All attributes-value-pairs or array items are separated by commas. Indentation and newline characters are not parsed in JSON as all syntax is given using the characters `""[]:'.',[32]`

2.7 Basic Concepts of Music Theory

Music itself is a form of art that combines either vocal or instrumental sounds, sometimes both to express an emotion or convey an idea [33]. But despite being able to define music it can be quite difficult even for artists itself to describe it and they struggle to put their perception into words. Famous Tenor and Saxophonist of the 1950s/1960s John Coltrane described Music as follows. "My music is the spiritual expression of what I am my faith, my knowledge, my being. When you begin to see the possibilities of music, your desire to do something really good for people, to help humanity free itself from its hang ups. I want to speak to their souls" [34].

2.7.1 Melodical Components of Music

Based on John Coltrane's quote, music could colloquially be described as something, "that is used to trigger human sense" [34]. To achieve this music consists of various elements that work together harmoniously to form a pleasant-sounding melody. Many factors are highly interdependent and correlate with each other.

- **Pitch and Melody:** Pitch is the word used to describe the highness or lowness of a musical sound. A series of Pitches, also described as scales, is used to create a Melody. Melodies can be derived from various scales for e.g., the traditional major and minor scales of music or even more unusual ones like the whole tone scale [35, p. 86]. Furthermore melodies can be described in two ways. Conjunct melodies are smooth and easy to sing or play. Disjunct melodies however represent the opposite. They are ragged or jumpy and difficult to sing or play.
- **Harmony and Chords:** Harmony can be described as the sound created when two or more Pitches are performed at the same time to form a chord. Chord is simply

the definition for three or more notes sounding at once. Harmony can be split up into two categories taking the sound of pitches into consideration. Harmony can sound Consonant, which means the pitches sound pleasant together or Dissonant, meaning the pitches sound unpleasant together [35, p. 94].

- **Rhythm:** Rhythm refers to the recurrence of notes and silences in time. A rhythmic pattern is formed by a series of notes and silence repeats. In addition to signify when notes are played, rhythm also defines how long notes are played and with what intensity. The difference in time creates varying note durations as well as different type of accents [36].
- **Texture:** Texture indicates the number of instruments or voices that are used to contribute to the density of the music. Texture can be split up into 3 types of Monophonies, Homophony and Polyphony. Monophony describes a single layer of sound e.g. a solo voice. Homophony is a melody with an accompaniment which can be a lead singer with a band or a solo singer and a guitar or piano. Polyphony is a form of texture which consists of two or more independent voices. One voice forms the melody and the other forms a support role. This could be for example a lead singer with a choir [37].
- **Timbre:** Timbre, also known as tone colour or tone quality, can be described as the specific tone or quality an instrument or a voice has. Timbre helps to distinguish instruments from each other when playing the same melody or simple notes. For example, a “C” Note on the Piano and a sung “C” have the same pitch but different sound quality which gets differentiated by timbre. Timbre usually can be described with adjectives used to describe color, temperature, or the human voice e.g., warm, cold, metallic, harsh, dry [35, p. 94].

2.7.2 Lyrics and Instruments

Lyrics are the definition for the linguistic part of a song. They usually consist of verses and choruses and can be implicit or explicit. More or less every song uses unique lyrics but nevertheless we can categorize lyrics by topics they address [38]. Statistics have shown that the most common themes are Growing Up, Friendship, Statements of Discontent, heartbreak or Death [39]. Another factor that must be taken into consideration when characterizing music are instruments. In the previous chapter we already talked about the fact that timbre is used to distinguish instruments from each other. The timbre focuses mainly

on the sound of the instruments to distinguish them. However, we can distinguish instruments not only by the way they sound but also by how the sound is produced in the first place. This is done with a classification of the instruments into 5 categories [40].

- Idiophones: Sound gets produced by the body of the instrument vibrating e.g., xylophones
- Membranophones: Sound produced by vibration of a tightly stretched membrane such as drums
- Chordophones: Sound produced by vibrating strings e.g., piano or cello
- Aerophones: Produce sound by vibrating columns of air such as the pipe organ or oboe
- Electrophones: Produce sound by electronic means e.g., synthesizers or electronic instruments.

2.7.3 Chronological Classification

2.7.4 Genres

One potential classification criterion would be the chronological categorization with the help of epochs. Due to the fact that modern music is diverse, epochs would be an insufficient solution. As a result genres were developed which include both chronological aspects and elements of classical music theory. Genres combine songs that sound the same or emit the same feelings, as well as songs that are based on the same characteristics of lyrics and instruments in songs. To further specify songs, genres again can be split into more homogeneous sub-genres [41].

Hip Hop

Although hiphop is often widely considered as a synonym for rap music it can rather be defined as somewhat of a cultural movement or a form of lifestyle that includes a music genre. Hiphop culture consists of various elements which are united under this umbrella term. Elements are djing/turntablism, rapping, beatboxing, breakdancing, and graffiti/visual art [42]. These typical street style elements or activities created a cultural revolution and shaped music styles, fashion, technology, art, and more [43]. Even to this day hiphop continues to develop new art forms that impact the lives of new and old generations. When it comes to the history of hiphop the culture has its origins in the 1970s where

it first appeared in African American ghettos. Because of this heritage, hip-hop still sees itself as street culture, and musicians who practice this form of music often have close ties to gangs, clans, and poverty [44]. Basic elements of hip-hop come from the genres of funk and soul, which are also influenced by African Americans. Especially because of the distinctive rhythms, these genres offered themselves very well as a basis. Through the development of hip-hop, however, it is also becoming more and more unique in type and form. The main characteristic of the music of hip-hop is the interaction between a rapper and a beat. Rarely are real instruments used for the music and beats are usually created by electronic instruments or samples which is, the use of already finished or existing sound/music recordings. Beats can be very diverse and can represent many moods, e.g., aggressive, relaxed, harsh, etc [42]. But the main focus in rap is usually on the rapper himself and even more on the lyrics. Rappers use rhythm, lyrics, and the timbre of their voice to express themselves. The Artists use their voice as an instrument and more importantly different voice pitches to adapt to the intention of the lyrics. Rappers are often measured by the scene on their so-called "flow", the ability how quickly and how smoothly they can perform chants without errors. The themes of hip-hop songs are often poverty, drugs, violence, struggles, or righteousness. Since many hip-hopers come from ethnic and racial minorities, they often use personal life experiences for their lyrics. In addition, many songs also want to convey a message that often revolves around necessary political changes, social justice, or grievances in society [45].

Jazz

Jazz is a musical style in which improvisation plays a central role. In many jazz performances, artists often play pieces they just made up from their heads and perform them on spot. Jazz has its historical roots in the early 19th century America more specifically in New Orleans. The city was an ideal breeding ground for jazz music because of given cultural diversity. The percentage of ethnic minorities was much higher in this city than elsewhere in America, which is why it was often called a melting pot of cultures [46]. The biggest influence on Jazz had the Afro-American culture as the music somewhat reflected the breaking away from previous rules of slavery and oppression. The improvisation of songs should act symbolically as a counter to the rules they had to deal with for a long time. Since its inception, jazz has gone through many different phases. The beginnings of jazz, from the 1910s to the 1920s, were characterized by small bands, often consisting of only a frontman and a few accompanists. The frontman often improvised pieces using a cornet or a trumpet, while the accompanists supported him with clarinets or trombones. In addition, often instruments like the banjo, piano, double bass, or drums were used to create a rhythm. In the 1930s and 1940s, the Swing and Big Band era emerged. Now, for the first time, singers appeared before big bands and bandleaders, and the clarinet was

largely replaced by the saxophone. Moreover, the jazz epicenter shifted from New Orleans further and further to New York [47]. The 1950s and 1960s then introduced laid-back cool jazz in contrast to the more fast-paced songs of the previous decades. Here a jazz quartet often played soothing and slow songs. With the introduction of electronic instruments in jazz from 1970 onwards, many subgenres were formed, such as jazz-rock. Until today, there are many forms of jazz, which are mainly oriented to the New Orleans origins and the 3 mentioned eras. Jazz has its main musical roots in the blues, but there are also elements of rock and classical music in it [48]. A distinctive rhythm is a key characteristic of jazz music, these are created mainly by "swinging" eighth notes. The so-called swinging is created by emphasizing one note of the eighth note pair while the second note is lighter and "swings" to the next note. Jazz is also very polyphonic, which means that many sounds are played simultaneously and as a result, various layers of harmony are laid over an initial basic melody. In addition to the use of classical European instruments such as the saxophone or trumpet, instruments such as drums, bass, keyboard, guitar, and trombone are often used. Some types of jazz also have front singers, but often pieces are played without vocals [49]. As already mentioned before the main characteristic of Jazz is the spirit of improvisation. This unifies almost all forms of jazz music. Jazz artists attach great importance to imprinting their own sound and style on the music, so they usually even play their own songs slightly modified or with a distinct style. This leads to the fact that thousands of jazz recordings can be found for the same song, but they all sound different. Finally, it is important to mention that jazz can also reflect many different emotions. Everything from pain to joy is possible. For many People of Colour, it resembles the feeling of freedom, because for them, as mentioned above, jazz represents a strong voice against suffering, oppression, and injustice [50].

Rock

The musical genre rock is a popular music genre that combines elements of rhythm and blues, jazz and country music while adding electric instruments. It originated as Rock 'n' Roll in the late 1940's and early 1950's and of course also is constantly changing and evolving. The basis of rock formed the music genres blues, gospel and country. Early Rock 'N' Roll came from cities like Memphis, Chicago, New Orleans or St. Louis. However, the genre spread very quickly throughout Western culture, and so it was mainly the British who liked the genre very much and who, in turn, took it to new heights [51]. British bands like the Beatles or the Rolling Stones emerged and became very popular in the USA as well. Rock dominated the music with its loud, dynamic, energetic and intense style for almost over 50 years until it was replaced by hip-hop as the most mainstream genre. One of the main characteristics of rock is the infectious beat and rhythm. The music was primarily designed for dancing and was a clear distinction from other music genres popular at the

time, such as jazz or swing. That's why from its start in the 50s the genre was very popular among young people as they felt they could break out of rigid traditions. Furthermore it is also important to mention that rock, like hip-hop later on, not only had an impact on music but also a cultural influence on the generations from the 1950s to the 1990s [51]. Rock especially supported the sense of rebellion and social justice in the western world of the 1960s and influenced clothing style, hair style and even attitude of its listeners for over 40 years. The older, more conservative generation at the time, which favored more quiet songs, rather detested rock. The energy already mentioned above is a unique characteristic that sets rock apart from many other genres. Rock music is often very wild, impulsive and driving. But what defines rock the most is the use of electric instruments and especially the use of the electric guitar [52]. The indispensable electric guitar makes rock probably the genre that is most influenced by a single instrument. Pioneers of rock like Elvis Presley, Jimi Hendrix or Chuck Berry experimented a lot with the electric guitar and used the unique sound to their advantage. Amplifiers allowed the artists to reach new melodic aspects and pitches that are not achievable with the use of pure acoustic instruments [52]. Rock music is typically performed only in bands with a lead singer. He usually plays an electric guitar himself and is accompanied either by other electric guitars, normal guitars, other electric instruments and a drum kit. Lyrical texts of rock are also very diverse due to its division into many subgenres. Lyrics may not be very profound other lyrics by artists like Bob Dylan, however, are considered comparable to fine poetry. Rock music subgenres are very different and vary greatly in terms of rhythm and tempo. For example, the music genre of heavy metal is almost incomparable with soft rock [53].

3 Implementation

3.1 Data Collection

In this section the approach and implementation of data collection for this project is examined. First, the basic requirements for the dataset are given and some existing datasets are evaluated. Then the resources used are presented and the general approach to data collection is explained, including authorization and how features and labels are collected. Lastly, the concrete implementation in Python is explained.

3.1.1 Requirements for the Dataset

Basic requirements the dataset should fulfill are

- **Includes Spotify song features**

Spotify provides a set of song features that were generated using their own models. The dataset should include these features, as they are needed to train the model.

- **Includes genre as label**

The dataset needs to include the genre of the track to use as a label for the classifier

- **Has sufficient sample size per genre**

In order to train the model well, a sufficient sample size is needed per genre. It was not known before collecting the data, how many samples are enough to get a decent result, which is why the dataset with the highest sample size is preferred, without compromising too much on the other criteria.

- **Song and Artist name**

The best way to filter out duplicates is to use the song and artist names. Spotify does provide a track id for each song, however, if a song is released twice (e.g. as a single and later in an album), these track ids will differ which will lead to a duplicate entry.

Additional fields are not going to be used in this analysis, but might still be collected in order to publish the dataset and enable others to use it for different applications.

3.1.2 Ressources and Approach

Spotify provides extensive documentation for developers on their developer website.[54] This includes a developer dashboard and the Web API documentation, which is the main resource for data collection from Spotify.

There are a number of pre-made datasets available online containing song data from Spotify, which were examined for this project. In conclusion, none of them were able to fulfill all of the requirements listed above, be it because of unclear methodology or an insufficient sample size. To get around this, a dataset was specifically created for this paper using the Spotify Web API.

The API offers a number of endpoints, which return JSON metadata directly from the Spotify Data Catalogue[29]. Requests are made via HTTP GET or POST methods. The API can be used by anyone, but authorization via the OAuth protocol is required to access data. To explore the API and find endpoints to use, Spotify provides a developer console, which can be used to send requests and see what kind of responses come back. This is not suitable for saving the data or making multiple requests programmatically, but is helpful for API exploration. As there is not one single endpoint that delivers all required fields, multiple queries that build on top of each other have to be made.

The approach began with using the API reference to get an overview over the endpoints and their responses. The specific endpoints that might return interesting data were queried using the Spotify Web Console to see a response with live data and which exact fields are returned. Once exploration was complete, a data collection workflow was implemented in Python 3, mainly using the libraries `http`, `json` and `requests`. The final result was saved as a CSV file to be used for data exploration and further processing.

Many API endpoints give the option to specify a country and language the results should be returned for. This was set to english and the United States wherever possible, to get a dataset able to be used in applications beyond this project.

3.1.3 Getting Features

In order to predict the genre of a track based on audio features, these features have to be requested for every track. Spotify provides an endpoint to get audio features for up to 99 tracks at a time. The typical request/response pattern for the audio-feature request endpoint is shown in figure 10.

Figure 10: Audio Feature Request

GET	https://api.spotify.com/v1/audio-features?ids={ids} <i>request audio features for id</i>
Parameter	
ids	comma seperated list of up to 99 song ids
Response	
application/json	
200	ok
<pre> 1 { 2 "audio_features": [3 { 4 "danceability": 0.677, 5 "energy": 0.638, 6 "key": 8, 7 "loudness": -8.631, 8 "mode": 1, 9 "speechiness": 0.333, 10 "acousticness": 0.589, 11 "instrumentalness": 0, 12 "liveness": 0.193, 13 "valence": 0.435, 14 "tempo": 82.810, 15 "type": "audio_features", 16 "id": "2e3Ea0o24lReQFR4FA7yXH", 17 "uri": "spotify:track:2e3Ea0o24lReQFR4FA7yXH", 18 "track_href": "https://api.spotify.com/v1/tracks/2e3Ea0o24lReQFR4FA7yXH", 19 "analysis_url": "https://api.spotify.com/v1/audio-analysis/2e3Ea0o24lReQFR4FA7yXH", 20 "duration_ms": 211497, 21 "time_signature": 4 22 }, 23 ... 24] 25 }</pre>	

With the exception of *type*, *id*, *uri*, *track_href* and *analysis_url*, all of the fields included in this response can be used as features in the dataset. However, this api call expects a track id, which needs to be retrieved using other api calls first. This could be a search endpoint, getting all tracks in a playlist, etc. Also, it does not give the track or artist names and doesn't include a genre.

3.1.4 Getting Track IDs and Labels

There is no simple endpoint that takes one or more track ids and returns a "genre" field in its response. The exploration of the API revealed two ways of getting the genre of a track.

The first way is using the artist of a track. Given a track id, the artists of the track and their corresponding ids can be requested by using the `/tracks/id` endpoint. Then, using the artist id, the genres that an artist is known for are returned, as shown in figure 11.

Figure 11: Artist Request


GET	https://api.spotify.com/v1/artists/{id} <i>request information about an artist by their id</i>
Parameter	
id	id of the artist
Response application/json	
200	ok
<pre> 1 { 2 "external_urls": { 3 "spotify": "link to resource.." 4 }, 5 "followers": { 6 "href": null, 7 "total": 15554811 8 }, 9 "genres": [10 "conscious hip hop", 11 "hip hop", 12 "north carolina hip hop", 13 "rap" 14], 15 "href": "link to resource...", 16 "id": "6l3HvQ5sa6mXTsMTB19rO5", 17 "images": [links to album covers...], 18 "name": "J. Cole", 19 "popularity": 89, 20 "type": "artist", 21 "uri": "spotify:artist:6l3Hv..." 22 }</pre>	

The exemplary API response shows a problem with this approach. One artist can be sorted into multiple genres. A given track might be associated with either of the artist's

genres, but the data does not show, which one exactly. Additionally a track might have multiple artists which further complicates this. Given these circumstances, this approach is problematic.

The second way is Spotify's "categories" feature. The app's search tab provides a number of categories that a user can browse to find new music in their preferred genre or style. In figure 12a an overview over some of the categories that are available in the Spotify app is shown. There are categories of multiple types like places and activities, but also for all major genres. In the app screenshot for example "Classical", "Jazz" or "Soul". These categories can be used to get tracks that belong in each specific category. When a user taps on one of the categories, playlists that contain tracks of the respective category are shown to the user (figure 12b).

Figure 12: Categories and Playlists in Spotify App



The API mirrors the app's behaviour and provides an endpoint to get a list of categories and their ids, one to get all playlists and playlist ids in a category, and one to get all tracks and track ids in a playlist. This chain of API calls is used to request every track in every playlist in a certain category.

3.1.5 Python Implementation

This section describes, how the previously explained API calls are used to request data from the API and save the complete dataset as a CSV file. As explaining the Python code in detail is beyond the scope of this paper, only the general approach to the application is explained here. A well documented Jupyter Notebook containing all of the code is contained in the appendix and should be referred to for further details (??).

The script was implemented in a way that lets anyone with a Spotify developer account run it to gather their own data. It supports filtering for certain genres to speed up the process and exporting the current data after each step to be able to pause the data collection and check the current results or continue later. The details of these functions are explained in this section using code snippets.

The Python *requests* library is used to execute the requests, as it supports the features needed to send the GET and POST requests needed without having to write complicated code. The *json* library is used for transforming Python dictionaries into json data, *math* for access to rounding functions, *os* for managing system paths and access to files, *http* to define retry and timeout behaviour and *csv* to finally transform the json data into a CSV file.

After importing all necessary libraries, the credentials are read and a method is defined for authenticating to the API. The http library is configured to minimize errors and exceptions due to late responses or failed requests. A filter is implemented to limit the categories the data should be retrieved for using category ids. Next, the four main steps of data collection are executed:

1. Getting categories
2. Getting playlists
3. Getting tracks
4. Getting audio features

For each step, a method was implemented, which takes data from the previous step and uses it to request further data from the API. Category ids are used to get a playlists category, playlist ids are used to get a playlists tracks and finally track ids are used to request audio features.

After going through all stages of data retrieval, the following JSON datastructure has been built:

```

1 {
2   "categories": [
3     {
4       "id": "hiphop",
5       "name": "Hip-Hop",
6       "playlists": [
7         {
8           "id": "37i9dQZF1DX0XUsuxWHRQd",
9           "name": "RapCaviar",
10          "tracks": [
11            {
12              "id": "2AaJeBEq3WLcfFWly8svDf",
13              "name": "By Your Side",
14              "album": {
15                "id": "2RrZgDND03MLu6pRJdTz5",
16                "name": "By Your Side"
17              },
18              "artists": [
19                {
20                  "id": "45TgXXqMDdF8BkjA83OM7z",
21                  "name": "Rod Wave"
22                }
23              ],
24              "features": {
25                "danceability": 0.649,
26                "energy": 0.508,
27                "key": 8,
28                "loudness": -10.232,
29                "mode": 1,
30                "speechiness": 0.0959,
31                "acousticness": 0.0345,
32                "instrumentalness": 3.59e-05,
33                "liveness": 0.0736,
34                "valence": 0.405,
35                "tempo": 157.975,
36                "duration_ms": 194051,
37                "time_signature": 4
38              }
39            },
40            ...

```

This JSON contains all the necessary fields needed for continuing with the CRISP DM process. To prepare the dataset for data analysis, the JSON data is transformed into a flat structure and stored in a CSV file. The final structure of the data is shown in table ??.

Tabelle hier einfügen

3.2 Data Understanding

In this chapter, the data itself is covered and explained. This starts at the origin of the data, meaning how and why spotify collects the data, followed by covering the features the

spotify algorithm uses and explaining how each of them works and what they describe. After these explanations, to get a deeper understanding, the different features are compared with the help different metrics.

3.2.1 The Streaming Service Spotify

Spotify is a music streaming service and has established itself as the Number one in this market. Since their launch in 2006, the company has gained over 365 million users, of which nearly 45 percent are subscribed to the chargeable premium service. This Premium service comes at a subscription cost of nowadays 12,99€ per month, giving the opportunity to listen to every song that is available at Spotify - this being ca. 70 million tracks from over 1,2 million Artists. [55]

Similar to other streaming services, Spotify also generates a lot of data. The question here is how to filter out data from music in order to personalize user accounts. The founders evaluated millions of songs and taught an AI to independently recognize patterns and certain structures within different genres. The results of the algorithm can be seen above all in the "Discovery Weekly Playlist". Users are presented with thirty songs that match their tastes here every Monday. At the moment, the Spotify algorithm, which to this day goes by the name Echo Nest, works because of three main components. [56]

First of all, it creates an individual Taste Profile for each Spotify user. There, it is recorded which artists, songs, albums, playlists and also podcast and audio books the respective user listens to. And also how long he listens to the music tracks, where he does it and how often. This data is then compared with the Taste Profiles of other users. For example, if another user has a large overlap in musical taste with your own playlists, you will be shown songs in the Discovery Weekly playlist that the other user already listens to regularly. Playlists with more followers end up in the Discovery Weekly folders more often than those with fewer followers. [56]

In addition, the algorithm analyzes the individual components of the music tracks. It breaks down each song available on Spotify by tempo, instruments, duration, highs and lows, timbre, and so on. Based on the large amount of data, the Spotify algorithm divides the songs into more than 1500 different genres so far. However, in addition to many regional differences in music genres, "nonsensical" genres keep coming out in the AI's analysis. While it might be able to imagine something under deep power pop punk, genres like "djent" are not really close to reality. [57]

This analysis step also includes the classification of songs into happy, sad, melancholic or other types of music, so the classification according to feelings. However, this feature

is also controversial among experts, because there is no actual definition of sad or happy music. That always depends on the individual listener and his mood. To support the AI at this point, it reads the titles of the playlists of the users. If a song often lands in playlists with titles like love songs or romantic music, the AI assumes that the song is to be rated as a love song. [57]

At last, the Echo Nest algorithm searches the entire web daily for blog posts, comments on websites, articles, YouTube videos, Facebook postings and comments and tweets, evaluates them, and uses them to create an opinion about a particular song. Thus, within a few days, the algorithm notices how a song is received overall. This quickly creates a specific rating for each track, which in turn influences the Discover function and other recommendation formats within the platform. [57]

3.2.2 Feature Analysis

In the course of the analysis carried out here, three different genres were selected in order to subsequently assign songs to the correct genre on the basis of the underlying data. The genres selected here are *rock*, *hiphop* and *jazz*. The selection was made under the original, subjective assumption that these genres differ particularly strongly in their characteristics and style. The assumption focused primarily on characteristics such as vocals, the general use of instruments, the instruments used, or the energy of the songs of the respective genre.

The important insight for further consideration of the analysis is here: Each genre that is considered, each song feature that is evaluated does not reflect reality, but only reflects the results of the Spotify algorithm. The analysis can, according to subjective opinion, assign a *hiphop* song to the category, or the genre *rock* without being wrong - since the assignment makes sense based on the Spotify evaluations. Since the categories in the selected genres are very accurate due to their popularity, the analysis is performed under the assumption that the two are congruent.

As already mentioned, every song has given thirteen Features, which break down his main attributes into numbers. To understand the data, it is crucial to get to know these features. For a better presentation, they are divided into four clusters below, Musical Standards, Mood, Properties and Context.

Musical Standards

This Cluster includes the features, that capture and reflect the standard properties of music. These features are *duration*, *key*, *mode* and *time signature*. The feature *duration*, as

the name suggests, holds the information of how long the Track is. This is measured in seconds. The feature *key* is measured with an Integer between -1 and 11 and holds information about the key the track is in. Integers map to pitches using standard Pitch Class notation, E.g., $0 = C$, $1 = C\sharp / D\flat$, $2 = D$, and so on. If no key was detected, the value is -1 . *mode* is an integer (either 0 or 1) and indicates the modality (major or minor) of a track. Major is represented by 1 and minor is 0 . At last, the feature *time signature* is measured in an integer between 3 and 7 . It is a notational convention to specify how many beats are in each bar (or measure). The time signature ranges from 3 to 7 indicating time signatures of " $3/4$ ", to " $7/4$ ". [58]

Mood

Following these features comes the cluster "Mood". This cluster lists those features that measure the emotive values in songs, i.e. whether the song encourages dancing, spreads a positive or rather negative mood, etc. This is captured with the features: *danceability*, *valence*, *energy* and *tempo*.

All of these features except *tempo* are measured in a float between 0 and 1 . *danceability* describes how suitable a track is for dancing based on a combination of musical elements including tempo, rhythm stability, beat strength, and overall regularity. A value of 0.0 is least danceable and 1.0 is most danceable. *valence* describes the musical positiveness conveyed by a track. Tracks with high valence sound more positive (e.g. happy, cheerful, euphoric), while tracks with low valence sound more negative (e.g. sad, depressed, angry). *energy* represents a perceptual measure of intensity and activity. Typically, energetic tracks feel fast, loud, and noisy. A value near 1.0 indicates high energy, while tracks near 0.0 indicates low Values in this Feature. At last, *tempo* holds the information about the overall estimated tempo of a track in beats per minute (BPM). In musical terminology, tempo is the speed or pace of a given piece and derives directly from the average beat duration. [58]

Properties

This next cluster includes all features that capture the audio characteristics of the tracks, such as *loudness*, *speechiness* and *instrumentalness*.

loudness is measured in a float number between -60 and 0 and measures the overall loudness of a track in decibels (dB). The values are averaged across the entire track and are useful for comparing relative loudness of tracks. *speechiness*, measured in a float number between 0 and 1 , detects the presence of spoken words in a track. The more exclusively speech-like the recording (e.g. talk show, audio book, poetry), the closer to 1.0 the attribute value. Values above 0.66 describe tracks that are probably made entirely of

spoken words. Values below 0.33 most likely represent music and other non-speech-like tracks. *instrumentalness*, measured in a float between 0 and 1 as well, plays the counterpart to *speechiness*. It predicts whether a track contains no vocals at all. "Ooh" and "aah" sounds are treated as instrumental in this context. The closer the *instrumentalness* value is to 1.0, the more likely the track contains no vocal content. Values above 0.5 are intended to represent instrumental tracks, but confidence is higher as the value approaches 1.0. [58]

Context


At last, this Cluster unions the Features *liveness* and *acousticness*. These capture everything around the Songs, for example if it is played in front of a live audience.

Both Features are measured in a float number between 0 and 1. *acousticness* is a confidence measure from 0.0 to 1.0 of whether the track is acoustic. 1.0 represents high confidence the track is acoustic. *liveness* detects the presence of an audience in the recording. Higher *liveness* values represent an increased probability that the track was performed live. A value above 0.8 provides strong likelihood that the track is live.

3.2.3 Overall Correlation

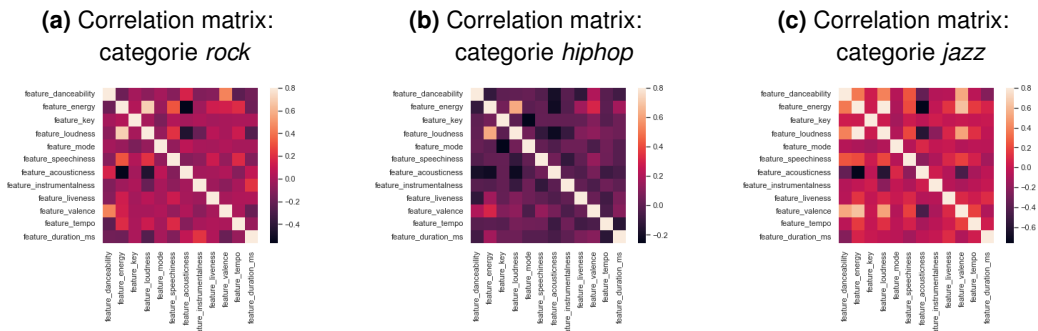
After these explanations, the characteristics of the features in the individual categories are evaluated. A correlation matrix is used for this purpose. This shows on a scale between 1 and -1 how features correlate with each other. Since every song uses every feature, a high correlation here does not mean that the features often appear together, but how similar the values per song are. To start with an overall look over the data, fig. 13 shows a genre independent matrix. This means that the features of all songs that could be gathered in the step of data collection were analyzed to create this figure. [58]

Figure 13: Correlation matrix: categorie-independent correlation of features



Generally, the features of the "Musical Standards" cluster can be ignored in this representation, as they do not measure but only reflect characteristics (e.g., time signature or key). It can be seen that *energy* and *loudness* have the highest correlation of the features, with a value of approx. 0.65. In addition, there is a comparatively high correlation between *valence* and "Danceability" of about 0.5. The correlation between *energy* and "*acousticness*" is particularly negative, with a value of approx. -0.7 . The same can be said for *acousticness* and *loudness*, with a correlation value of about -0.5 . Outside of these extreme values, the remaining correlation values settle between -0.2 and 0.2 . Worth mentioning here are the correlations between *instrumentalness* and *loudness* at approx. -0.3 as well as *instrumentalness* and *Valence* also at approx. -0.3 . Furthermore, in the positive *speechiness* and *energy* at a value of approx. 0.3 .

Figure 14: Correlations inside categories



These correlations above now show the matrix on a category-specific basis. It can be seen that the patterns described above also appear there. However, the average correlation changes depending on the category. For example, hip-hop has the same characteristics as the independent matrix category but shows fundamentally more negative correlations (cf. Fig. 14).

3.2.4 Comparison of Music Theory to Audio Features

In the final step of Data Understanding, the music theory described in the basic section is compared to the data used in the analysis. This is necessary to develop a deeper understanding of the data. As mentioned at the beginning of this chapter, theory and Spotify are not always necessarily congruent, which is why a clear delineation is important here. For this purpose, the properties of the genres in the music theory are tried to be compared with the help of the features (e.g., "You can't dance well to *rock*" - What is the overall value for Danceability?). To do this comparison, three features are chosen and

discussed more deeply. Subsequently, a better picture can be obtained with the help of further representations.

Starting with rock, one feature can be extracted here: The songs are mostly loud. *loudness* captures this property very well. It measures the actual, average loudness of a song in decibels, with most tracks falling in the -60 to 0 range, where -60 is very quiet and 0 is very high. Overall, across all three categories, this feature has an average of -8.2544 , with a minimum value of -32.06 and a maximum value of 2.044 .

Figure 15: *loudness* compared in the different categories

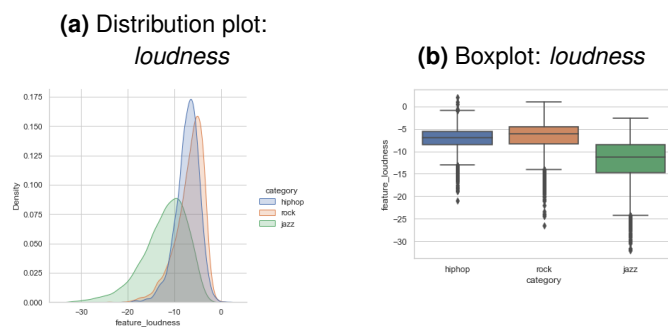
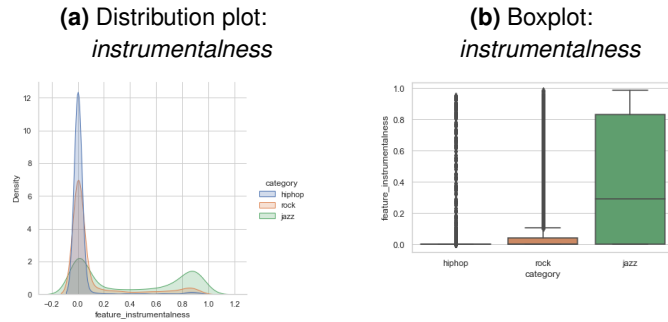


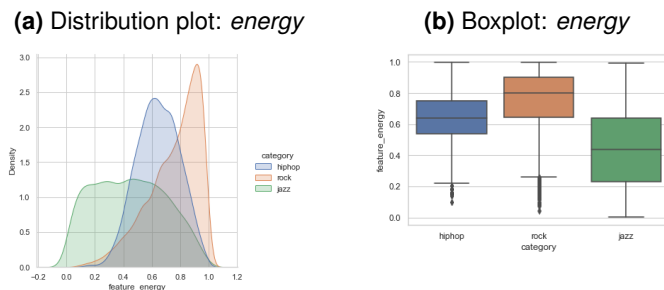
Figure 15 shows a so-called distribution plot. The X-axis contains the scale of the feature used (in this case *loudness*). The Y-axis reflects the density of the respective category, meaning the more songs are found in the respective value, the higher the density. As can be seen in the figure, all categories offer a very wide range of volume levels. *jazz* spreads its songs over an area from -32.06 (smallest value overall) to -2.695 but locates most songs at about -10 . *hiphop* and *rock* go less wide but can accommodate more songs in the upper quarter of the scale. *rock* places its point furthest to the right of the scale. Thus, it can be confirmed that *rock* can generally be classified as loud, but also shows as the loudest of the three categories. However, it must also be emphasized here that *hiphop* comes shortly after. In addition, *jazz* is not as loud as the other two categories, but it serves a much wider range.

The next feature that can be scrutinized here is the large occurrence of instruments in *jazz*. This can be checked in the Spotify data via the *instrumentalness* feature. It measures the sole presence of instruments in songs with values between 0 and 1. The closer the value approaches 1, the higher the probability that there are no vocals or other elements in the song. The overall average for this feature is 0.174. The lowest value overall is 0.0 while the highest goes up to 0.989.

Figure 16: *instrumentalness* compared in the different categories

For a better understanding, Figure 16 also shows a distribution plot, but adapted to the current feature. It shows that both hiphop and *rock* are almost exclusively found in the 0.0 to 0.1 range. This means that songs in these categories almost only contain vocals. However, slight increases in the range between 0.8 and 1.0 can also be seen. *jazz*, on the other hand, has two equally strong expressions. One in the range around 0.0, the other between 0.8 and 1.0. Since *jazz* often contains only musical tracks without any vocals, this reflects the reality quite well. Here again, *rock* and hiphop are similar in their high points. In addition, *jazz* much differs from the two above. Again, it does not have such a high peak, but covers the scale much more broadly.

At last, a comparison of the categories in the *energy* feature. Typically, energetic tracks feel fast, loud, and noisy. The higher the value approaches 1, the more energetic it is. Basically, it is to be noted that there is a high average value of 0.65 overall. This is clearer for *rock* with an average value of 0.7531.


Figure 17: *energy* compared in the different categories

If we now look at the data in a distribution plot (17), we see clear differences between the categories. *jazz* spreads over the whole graph but has no high point. This can be attributed to the different types of *jazz* (slow *jazz*, free *jazz*), which means that different subtypes have different levels of energy. Surprisingly, *hiphop* is also broadly positioned


and sets the high point at about 0.61. *hiphop* also has many facets as a genre, which is why this expression also makes sense. The curve of *rock* builds up early but rises very steeply in the last third of the graph and marks its high point at about 0.9.

After these three sample features, some more distributions and boxplots are added to better explain the characteristics of the individual genres.


(a) Distribution plot:
Speechiness




(b) Boxplot: *Speechiness*




(c) Distribution plot:
danceability



(d) Boxplot: *danceability*



(e) Distribution plot:
acousticness



(f) Boxplot: *acousticness*





Figure 18: Different features visualized in distribution- and boxplots

3.3 Data Preparation

In the following section the data is prepared and cleaned up for training the model.

Data preparation is done in Python using multiple libraries. The same libraries are also used in the modeling and evaluation stages of this project. *pandas* is used for storing and manipulating data in dataframes. *sklearn* is used for splitting data, creating validation curves, transforming data and training the gradient boosting model. *matplotlib* and *seaborn* are used for data visualization and plotting. Additionally *numpy* is used for other operations on the data.

As with data collection, the complete code with documentation is attached in the appendix.

3.3.1 Import and Cleaning

The CSV file is first imported and converted into a pandas dataframe.

```
1 df_import = pd.read_csv(os.path.join('..', 'data_collection', 'final_result.csv'))
```

Next, duplicates are removed. If there are two or more datapoints, which have the same value in their artist and name column, only the first one is kept and all subsequent entries are removed. This deduplication could also be done by using track ids. The drawback of this method is, that many artists release their tracks multiple times, e.g. as a single and later in an album. These duplicates would not be caught using the id, as different releases have different id values. As the artist and title doesn't change, almost all duplicates are caught using artist and name.

```
1 df_dedup = df_import.drop_duplicates(subset=['categories.playlists.tracks.artists',  
      , 'categories.playlists.tracks.name'])
```

Next, all genres that are not to be used for training the model are filtered out.

```

1 genre_filter = ['hiphop', 'jazz', 'rock']
2 df_filtered = df_dedup[df_dedup['categories.id'].isin(genre_filter)]

```

Then columns that are not needed for training are removed, including category name, all playlist information and all track and album information. The category id is kept, as it will serve as the label. The remaining columns are renamed, as the long JSON tree names are no longer needed. The field *category.id* is renamed to *category* and each audio feature is renamed for example the danceability column is now called *feature_danceability*. With unnecessary columns removed, a check is done to show any remaining null values.

```

1 df.isnull().sum()

```

In this dataset, there are no null values present.

The *GradientBoostingClassifier* used for modeling only supports integer values as label input. The category data is therefore encoded with integers using a custom function *encode_target*. It takes a dataframe and the label column's name, then adds a "target" column with corresponding integer mappings.

```

1 def encode_target(df, target_column):
2
3     df_mod = df.copy()
4     map_to_int = {name: n for n, name in enumerate(df_mod["category"].unique())}
5     df_mod["target"] = df_mod[target_column].replace(map_to_int)
6
7     return (df_mod)
8
9 df_target = encode_target(df, "category")

```

The head of the resulting dataframe is shown in table ???. The category to target integer mapping is shown in table 3.3.1.

Table 9: Dataframe after cleanup

Hier Category/target mapping table

Looking at the number of samples per category (tabel ?? reveals, that the dataset is very uneven, with more than half of the samples being in the category *rock*.

Hier samples per category

As explained in section 2.4.1 and 2.4.2 methods like standardization and PCA can have a positive impact on the models accuracy. Before finding the best model using hyperparameter tuning, the best form of input data is evaluated by transforming the data in different

ways and training a Gradient Boosting model using the default parameters specified in sklearn. This way, the best form of input data can be found without the overhead of resource intensive grid search.

To be able to easily train models using different forms of input data, a method *eval_prep* was created, which takes a dataframe, a list of all feature column names and the label column name. It then splits the dataframe into train and test sets, trains a *GradientBoostingClassifier* and returns a simple accuracy score. This process is explained in depth in the modeling section.

First, the regular dataframe is used as input data, resulting in an accuracy score of 0.8505.

```
1 score = eval_prep(df, features, "target")
```

3.3.2 Data Transformation

Next, mean removal, variance scaling and standardisation are attempted, by using the *StandardScaler* class from sklearn to transform the data. The code for standardization is shown below. Mean removal and variance scaling use the same class with additional parameters.

```
1 X_s = df[features]
2 y = df["target"]
3
4 X_s = StandardScaler().fit_transform(X_s)
5
6 df_s = pd.DataFrame(data=X_s)
7 df_s.insert(0, "target", y)
8 df_s.columns = ["target"] + features
9
10 score = eval_prep(df_s, features, "target")
```

This results in the following scores, beating the untransformed data in all cases.

- Mean Removal: 0.8509
- Variance Scaling: 0.8520
- Standardization: 0.8523

3.3.3 Principal Component Analysis


Next PCA is attempted using sklearn's *PCA* class. It is able to take any dataset and reduce its dimensionality to any number of dimensions. The dimensionality is reduced to all dimensions between 2 and 13 to see, which dimensionality results in the best score. PCA was done both on the raw dataset and in combination with standardized data (both before and afterwards).

PCA does not yield good results with this dataset. The score for all of the combination is always worse than using untransformed data. The PCA score was best when keeping 13 dimensions, declining further when removing more dimensions. These are the scores using each combination with PCA and keeping 13 dimensions:

- PCA only with 13 components: 0.8352
- Standardization after PCA with 13 components: 0.8348
- Standardization before PCA with 13 components: 0.8277

Figure 19 shows the test accuracy using PCA on the untransformed data for each number of components.

Figure 19: Test accuracy for each dimensionality after PCA



As the results using PCA do not improve the score, the code is not shown here. Testing data transformation and dimension reduction with this dataset and a default *GradientBoostingClassifier* shows, that the model benefits slightly from standardized data, improving the overall score compared to the raw dataset by 0.19%. Therefore, standardized data is used going forward.

3.4 Modeling

3.4.1 The Classifier

For training the Gradient Boosting Algorithm, the *GradientBoostingClassifier* class from the scikit learn Python library is used. It is part of a group of classes offering different ensemble methods. As explained in section 2.3, ensemble methods combine the predictions of several weak learners to generate a more robust model and reduce overfitting issues. sklearn supports averaging methods like Bagging and Random Forests, which take the average of each learners prediction as their output. It also supports Boosting Methods like AdaBoost or Gradient Boosting, which build base estimators sequentially, improving the output each time. sklearn also provides a classifier and regressor model for each method.

The Gradient Boosting Classifier supports binary and multi-class classification and uses 20 hyperparameters to control the size of each regression tree, the number of trees, the learning rate and many more. As explaining and tuning all 20 hyperparameters of the classifier would be beyond the scope of this paper, three important parameters are explained and optimized.

- **learning_rate**

As explained in section 2.3.1, the learning rate is used to control how much each tree contributes to the result, by multiplying it with the output values of the previous tree. The default value here is 0.1

- **n_estimators**

The number of weak learners (here regression trees) to be used while boosting. The default is 100.

- **max_depth**

The maximum depth of each regression tree. This also impacts its number of nodes. The default value is 3.

3.4.2 Validation Curves

During hyperparameter tuning, grid search is used to find the optimal value for each parameter. As this is very resource intensive, a sensible range of values to search in must be found beforehand. Validation curves can be used to see, how a model's score changes

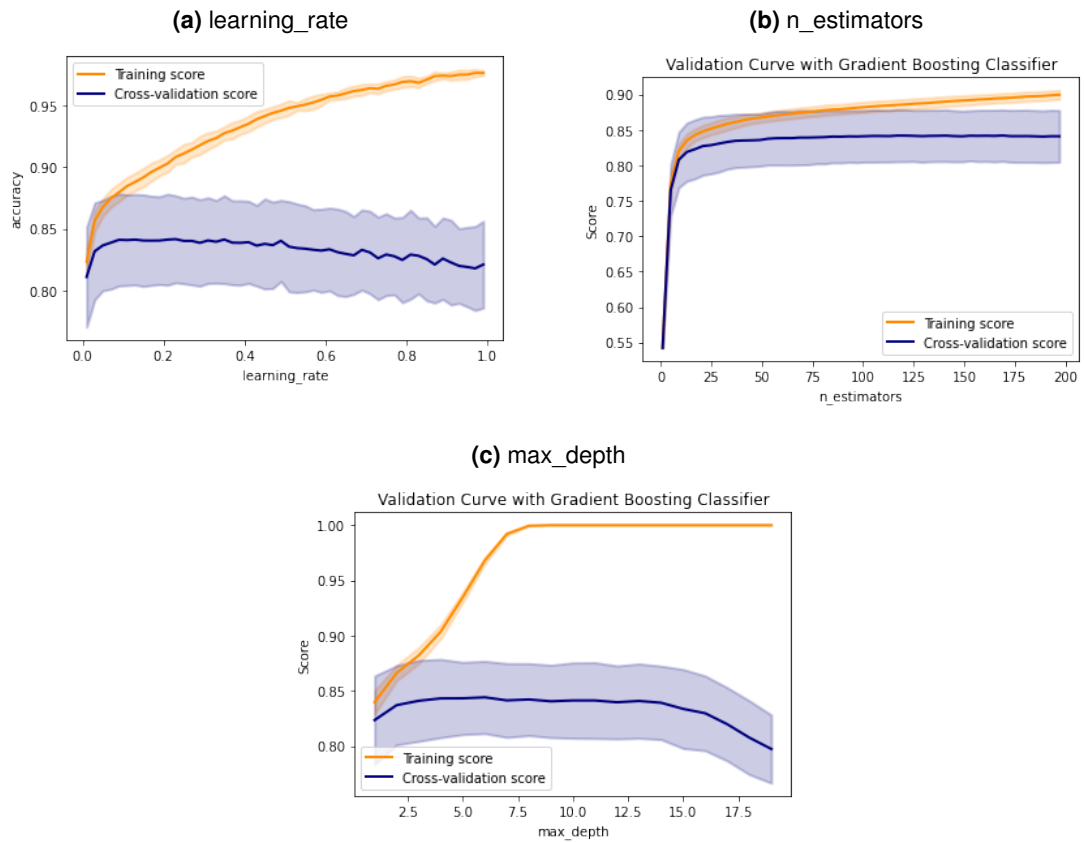
when modulating a single hyperparameter. Here, an important distinction between the training and cross-validation score has to be made. The training score is the accuracy the model achieves, when being scored on the same data it was trained on. A very high training score could be a sign of overfitting. The cross-validation score (cv score) is the mean accuracy achieved during 5-fold cross validation and is a good measure of the actual performance of the model.

Figure 20 shows the curves for each hyperparameter. The shaded areas around the graph are the standard deviation, while the graph itself represents the mean score between 5 folds.

Figure 20a shows the best cv score with a learning rate between 0 and 0.2, declining as it approaches one. The training score rises as the learning rate increases, which indicates, that the model becomes less generalizable and more overfitted as the learning rate increases. Because of these results, only values between 0.01 and 0.25 are considered during hyperparameter tuning.

The cv score stays relatively flat as the number of estimators goes beyond 25 (figure 20). The training score increases slightly. This indicates, that a training score close to, but above 25 might be optimal. As the results are not very clear, a range between 50 and 200 is used for optimization.

For *max_depth*, the cv score stays relatively flat between 2 and 13 (figure 20). The training score, on the other hand, increases rapidly from 0 to 7.5, where it reaches its maximum. Because of this, a range between 2 and 6 is used for hyperparameter tuning.

Figure 20: Validation Curves modulattng 3 hyperparameters

3.4.3 Hyperparameter Tuning and Fitting a Model

The parameter grid is given as a Python dictionary containing the parameter name as keys and a list of values for each key. All undefined hyperparameters are kept at default. The grid used for training with the value ranges found above looks like this:

```

1 param_grid = {
2     "learning_rate": [0.01, 0.04, 0.07, 0.1, 0.13, 0.16, 0.19, 0.22, 0.25],
3     "n_estimators": [50, 75, 100, 125, 150, 175, 200],
4     "max_depth": [2, 3, 4, 5, 6]
5 }

```

The grid and a *GradientBoostingClassifier* object are passed into a *GridSearchCV* object which executes the grid search using 5-fold CV.

```

1 gbc = GradientBoostingClassifier(random_state=45)
2 search = GridSearchCV(gbc, param_grid,
3     n_jobs=-1,
4     error_score="raise",
5     verbose=1)

```

The search object's fit method combines cross-validation, hyperparameter tuning and classifier fitting to find the best possible combination of hyperparameters to the estimator for the given dataset. This is done in the following way:

1. An estimator is created using the first combination of hyperparameters.
2. The training set is split into five folds and each fold is used once for calculating accuracy, with the other four used for training. The average of the five scores is used as the final score for this specific set of parameters.
3. The score is saved together with the parameters
4. These steps are repeated with every possible combination of parameters from the parameter grid. This results in a score for every set of parameters.
5. As every model was fitted using the same method, it is clear that the model with the highest score is using the best hyperparameters. A new model is trained without cross-validation using all available training data, to create the final model.

The complete standardized dataset is split in 80% train and 20% test data. The samples are shuffled before splitting to ensure an even distribution of data. Then features and labels are separated for input into the classifier.

```
1 train, test = train_test_split(df_s, test_size=0.2, random_state=45, shuffle=True)
2
3 X_train = train[features]
4 X_test = test[features]
5 y_train = train["target"]
6 y_test = test["target"]
```

Then hyperparameter tuning and estimator fitting is started using the training features and labels.

```
1 search.fit(X_train, y_train)
```

The parameter grid above has 315 possible combinations. As each is fitted on five sets of folds, there are 1575 models fitted in total during this process.

The resulting model with the best cross-validation score used the following parameters:

- learning_rate: 0.1
- n_estimators: 100
- max_depth: 3

Calculating the accuracy of the resulting model using the test set resulted in an accuracy score of 0.8591.

A weakpoint of the parameter grid used is, that the values for *n_estimators* are spread by increments of 25, leaving possible room for improvement. To check, if the model could be improved, another parameter grid was used:

```
1 param_grid = {  
2     "learning_rate": [0.01, 0.04, 0.07, 0.1, 0.13, 0.16, 0.19, 0.22, 0.25],  
3     "n_estimators": [97, 98, 99, 100, 101, 102, 103],  
4     "max_depth": [2, 3, 4, 5, 6]  
5 }
```

This did not deliver a better result.

Even though a broad spectrum of values was given in the search grid, the optimal values are equal to the default parameters for the estimator set by *sklearn*. This might suggest, that the developers set the default values using a similar dataset to the one used in this project.

Although many combinations were evaluated, there is still a possibility, that this is a local maximum and a set of hyperparameters that is far different from the ones tested here are actually optimal.

3.5 Evaluation

After the completion of the modelling phase, the next step is to evaluate the results. The evaluation is split into the evaluation of the model itself followed by the evaluation of the overall process.

3.5.1 Gradient Boosting Evaluation

In this section the gradient boosting algorithm will be observed in more detail by the most common metrics. Additionally, it also will be compared with the much simpler classification tree algorithm to highlight commonalities and different approaches between the two algorithms and create an all-encompassing picture for Gradient Boosting.

The overall accuracy measured is around 86%. This result is already positive as the Gradient Boosting Algorithm outperforms the classification tree, which achieves an also respectable 80% accuracy. However, the accuracy score can only be used as a fundamental

basis for evaluating the overall performance with many unknowns that need to be worked through.

The first step of a deeper analysis is to create a confusion matrix for all categories. The confusion matrix is an approach of visualizing the performance by clustering the the output. The x_1 -axis represents the predicted values for the classes while the x_2 -axis stands of actual (correct) values [59, p.235]. For this analysis three numbered classes 1 to 1 are necessary that typify the categories hiphop, rock and jazz in the stated order. The confusion matrix reveals four combinations of predicted and actual values (10). For Gradient Boosting the confusion matrix is displayed in table 11.

Table 10: Confusion matrix

predicted actual	class 1	class 2
class 1	true positive	false negative
class 2	false positive	true negative

Table 11: Gradient boosting confusion matrix

predicted actual	hiphop	rock	jazz	all
hiohop	438	68	19	525
rock	41	1307	122	1470
jazz	30	115	536	681
all	509	1490	677	2676

Visually, the highest misclassification takes place between classes *jazz* and *rock* while the lowest missclassification occurs between *hiphop* and *jazz*. This result however has little significance, since the number of *rock* tracks, with a total of over 1400, is significantly larger than the amount of both *hiphop* and *jazz* tracks. An evaluation based on absolute numbers would lead to misleading results partly due to the imbalance of the dataset.

A better approach is to look at the following metrics which can be derived form the confusion matrix for every category [59, p.235] [60, p.862](12).

$$\begin{aligned}
 \text{precision} &= \frac{\text{true positive}}{\text{true positive} + \text{false positive}} \\
 \text{recall} &= \frac{\text{true positive}}{\text{true positive} + \text{false negative}} \\
 f1 - \text{score} &= 2 * \frac{\text{precision} * \text{recall}}{\text{precision} + \text{recall}}
 \end{aligned}$$

Table 12: Gradient boosting classification report

classes	precision	recall	f1-score	support
hiphop	0.86	0.83	0.85	525
rock	0.88	0.89	0.88	1470
jazz	0.84	0.79	0.79	681

Precision and recall are both performance metrics with different objectives. While precision is a measure of how many of the predicted elements for a class were correct, recall measures how many elements of a category were detected. It can be observed that the precision is high across all categories with *rock* being classified best with 12% false-positive predictions while the false-positive rate for *jazz* was worst with over 21%. Interestingly the results for recall are very similar with the model performing best for the category *rock* with an recall of 0.89 and again worst for *jazz* with 0.79. *Hip-hop* is for both metrics in between of both extremas at around 0.85.

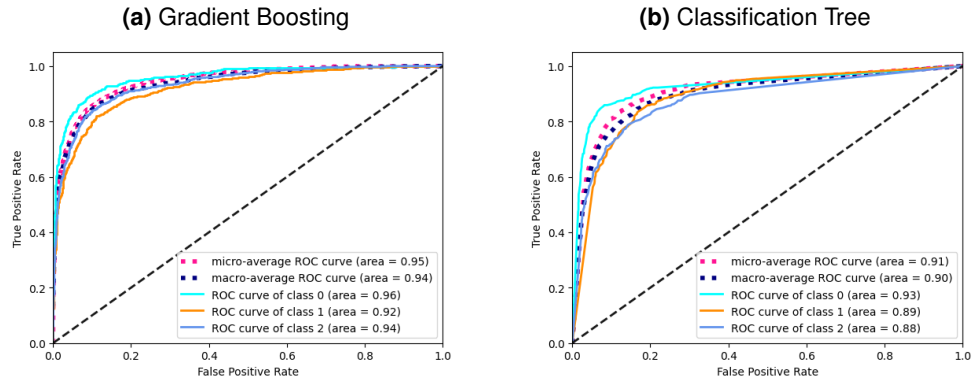
The f1-score is a combination out of precision and recall and an attempt of capturing both metrics in a single value. Therefore the results are not surprising. *Jazz* performed worst with only 0.79 while *rock* was classified best with 0.88 according to the f1-score.

In conclusion all scores can be evaluated as positive outputs without any negative and unexplainable abnormalities. Also in comparison to decision trees a clear improvement of the results can be seen. The reason that *jazz* is ranked worst while *rock* reaches the highest values may be due to several reasons. One possible explanation can be found by reviewing the findings from the data understanding chapter. It is recognizable that the value-ranges for the features of *jazz* were significantly more distributed than for both *rock* and *hip-hop* and often without any high points. In addition, the value ranges of *rock* were often different from those of hip-hop and jazz, which simplifies the respective classification.

Another measure to evaluate a model's performance is to plot its Receiver Operating Characteristic (ROC)s. The ROC is a plot of the true-positive rate for the x_1 -axis and the false-positive rate on the x_2 -axis for every possible threshold. The benefit of ROC is that it plots the misclassification for every threshold while other metrics rely on a single threshold. Often there are also requirements for models in which the overall accuracy plays a subordinate role, such as that all samples of class x must be detected. The ROC helps to visualize such problems to find optimal solutions. A model whose results are close to the diagonal classifies data worse than a model whose curve is as close as possible to the point (0/1) in the coordinate system. This performance is often measured by means of the Area Under the Curve (AUC), which is derived directly from the ROC. With help of


the AUC comparability between models is possible. The ROC and AUC for the project are shown in figure 21 [60, p.862f] [61].

Figure 21: Comparison of ROC and AUC



From the following matrix it can be seen that the model works well for all classes. Both ROC and the associated AUC are convincing in all cases. What is interesting is the fact that *rock* has the worst AUC while being classified best according to the previous metrics. *Hiphop*, on the other hand, has the significantly best AUC followed by *jazz*. An explanation for this result is complicated, with a possible reason again being the imbalance of the dataset. The model is focussed classifying *rock* correctly as it constitutes to a large part of the dataset. Therefore the optimum for both *jazz* and *hiphop* are exchanged for an overall optimum. It is noted that a well founded explanation would require more in-depth analysis. Regardless, the result is worth including in the analysis and leaves room for further research.

In addition, a very different evaluation can be performed. It is also interesting to take a closer look on the input data with help of the so-called feature importance. The feature importance is a measure of how much impact a feature had for the classification of the dataset and is presented in the following figure 22 for both gradient boosting and decision trees [62].

Figure 22: Comparison of feature importance

It is clearly visible that both gradient boosting and decision trees relied on similar features to similar extents for the classification task with minor differences for features such as danceability. Noticeable is that gradient boosting is heavily focused on a few features while the classification tree makes more use out of features like *duration*, *valence* and *mode*. One possibility for the difference in weighting could be the overall construction of the algorithms. Gradient boosting consists out of multiple small weak learners while the classification tree algorithm forms a widely branched and deep tree. For feature importance an analogy can again be made to data understanding and the music theory. The features that theoretically distinguish the genres from another play the most important role for the modeling while musical standards and more subjective features only play a subordinate role. This result is positive both for theory and the machine learning model. The result states that music theory can be confirmed by real-world examples while giving credibility and plausibility to the model.

3.5.2 Process Evaluation

The evaluation of the overall process is complicated as there are many unknowns starting with the data itself. The data used is preprocessed by Spotify without detailed definition and theoretical basis. The features are furthermore heavily subjective which adds another layer of complexity. Also the data collection is non-trivial and error-prone. Both the collection process itself and the collection approach would have to be revised for a productive use case. On the other hand, the data quality is decent with high correlation to the music theory. In addition the dataset convinces with its uniqueness and novelty.

The model in combination with the preprocessing and evaluation also leaves room for further research. The use of only two very similar machine learning algorithms allows no comparison to completely different approaches based on other algorithms and concepts. It

can therefore not be ruled out that even better results can be achieved. The pre-processing is intensive with many constallations tested. A possible improvement would be to apply the hyperparametric optimization not only on the standardized data but on all approaches to ensure the overall best result. In addition, there are further hyperparameter constellations which were not tested. The evaluation may also have reached incorrect conclusions and assumptions for a variety of reasons.

The overall highest risk is the fact that all participants of the group are not familiar with the Machine Learning and lack in experience. For this reason, the result can be considered promising.

4 Conclusion

The project had the purpose to classify a dataset with the help of a machine learning algorithm. With no restrictions on the choice of dataset and algorithm given, a classification of genres based on Spotify Music data was selected as a project task. After an extended search, no sufficient existing dataset was found and thus a new dataset had to be created. Even though the data collection was complex at times, it proofed to be a good decision and in convincing with its novelty. After testing and debating over various algorithms, gradient boosting was selected as the machine learning algorithm. Additionally, a classification tree was implemented as a reference point.

The overall result was successful. The gradient boosting model recieved an accuracy fo around 86% with additional in-dept evaluation metrics supporting the positive performance. It is worth mentioning that despite high efforts of optimization only standardization has lead to a better classification result as other models did not lead to an improvement. This indicates that either the models were implemented incorrectly or that the overall quality of the dataset was excellent from the beginning. The implementation itself, despite all the positives, still leaves room for additional research. It is furthermore very pleasing to see that the model relies similar features that are also used in music theory for differentiation between genres and allows affirmatory assumptions to be made on Spotifys data quality, the gathered dataset and the implementation of the gradient boosting algorithm.

Also from a personal perspective there are several lessons learned during this project. The first lesson is, that despite the choice of a comparatively simple machine learning algorithm in gradient boosting, limitations of the project framework were reached. Both in terms of scope and size of the paper and from an implementations perspective, as multiple optimization possibilities and design choices were discovered and required further

attention. The project was not straight-forward with multiple modifications and additions to be made along the way. However, all changes should in no way be seen as negative, as they have led to the desired result and are the result of a steep learning process. No participant of this group had prior experiences in machine learning. New discoveries had to be gathered and processed to form conclusions and drive the project forward. It is remarkable how uncomplicated the creation of simple models is with the help of machine learning libraries. Nevertheless, the final classification model is convincing and meets the predefined objectives. The result is a team effort and was only possible through good collaboration.

In conclusion, Spotify is just one example of many modern applications that use big data analytics to capture and attempt to reflect the often subjective world in numbers. What is particularly interesting here is how this is attempted, and to what extent it works. The success of the company, or the widespread use of the app, show how well the Echo Nest works. Still, like the model shown here, the deciding attributes are the ones closer to numbers, like Acousticness or Speechiness. Attributes such as Valence, which attempts to measure mood, plays a small role in genre assignment. However, this can, and will, change in the future. The decoding of the personal feeling of the user is the next step to more, and above all more decisive data and in the broader sense of a better customer relationship. So we can look forward with great interest to the further development of such use cases.

In conclusion, Spotify is just one example of many modern applications that utilize big data analytics to attempt to capture and reflect the oftentimes subjective worlds in numbers. Particularly interesting is the fact to what extent music, which is believed to be subjective, can be processed by Spotify to give recommendations, cluster songs in genres or feelings.

Appendix

Appendix 1: Beispielanhang

Dieser Abschnitt dient nur dazu zu demonstrieren, wie ein Anhang aufgebaut sein kann.

Appendix 1.1: Weitere Gliederungsebene

Auch eine zweite Gliederungsebene ist möglich.

Appendix 2: Bilder

Auch mit Bildern. Diese tauchen nicht im Abbildungsverzeichnis auf.

Figure 23: Beispielbild

Name	Änderungsdatum	Typ	Größe
 abbildungen	29.08.2013 01:25	Dateiordner	
 kapitel	29.08.2013 00:55	Dateiordner	
 literatur	31.08.2013 18:17	Dateiordner	
 skripte	01.09.2013 00:10	Dateiordner	
 compile.bat	31.08.2013 20:11	Windows-Batchda...	1 KB
 thesis_main.tex	01.09.2013 00:25	LaTeX Document	5 KB

Bibliography

- [1] M. Tanwar, R. Duggal, and S. K. Khatri, 'Unravelling unstructured data: A wealth of information in big data,' in *2015 4th International Conference on Reliability, Info-com Technologies and Optimization (ICRITO) (Trends and Future Directions)*, IEEE, 2015-09. DOI: 10.1109/icrito.2015.7359270.
- [2] A. Meier, 'Rundgang big data analytics – hard & soft data mining,' in *Big Data Analytics*, Springer Fachmedien Wiesbaden, 2021, pp. 3–23. DOI: 10.1007/978-3-658-32236-6_1.
- [3] 'Quick guide: Machine learning im maschinen und anlagenbau,' VDMA Software und Digitalisierung, 2018.
- [4] G. Paaß and D. Hecker, *Künstliche Intelligenz*. Springer Fachmedien Wiesbaden, 2020. DOI: 10.1007/978-3-658-30211-5.
- [5] S. Schacht and C. Lanquillon, *Blockchain und maschinelles Lernen: Wie das maschinelle Lernen und die Distributed-Ledger-Technologie voneinander profitieren*. Springer Berlin Heidelberg, 2019, ISBN: 9783662604083. [Online]. Available: https://books.google.de/books?id=nL%5C_ADwAAQBAJ.
- [6] G. James, D. Witten, T. Hastie, and R. Tibshirani, *An Introduction to Statistical Learning*, 2nd ed. Springer US, 2021. DOI: 10.1007/978-1-0716-1418-1.
- [7] R. J. Lewis, 'An introduction to classification and regression tree (cart) analysis,' in *Annual meeting of the society for academic emergency medicine in San Francisco, California*, Citeseer, vol. 14, 2000.
- [8] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay, *Scikit-learn: Machine learning in Python*, Feature importances with a forest of trees, <https://scikit-learn.org/stable/modules/tree.html>, Accessed: 2022-01-29, 2011.
- [9] S. Tangirala, 'Evaluating the impact of gini index and information gain on classification using decision tree classifier algorithm,' *International Journal of Advanced Computer Science and Applications*, vol. 11, no. 2, pp. 612–619, 2020.
- [10] A. Natekin and A. Knoll, 'Gradient boosting machines, a tutorial,' *Frontiers in Neuro-robotics*, vol. 7, 2013. DOI: 10.3389/fnbot.2013.00021.
- [11] T. Parr and J. Howard, *How to explain gradient boosting*, Gradient Boosting: Distance to target, <https://explained.ai/gradient-boosting/L2-loss.html>, Accessed: 2022-01-29, 2022.
- [12] P. Bühlmann, 'Bagging, boosting and ensemble methods,' Berlin, Papers 2004,31, 2004. [Online]. Available: <http://hdl.handle.net/10419/22204>.
- [13] T. Parr and J. Howard, *How to explain gradient boosting*, Gradient boosting: Heading in the right direction, <https://explained.ai/gradient-boosting/L1-loss.html>, Accessed: 2022-01-29, 2022.

-
- [14] J. H. Friedman, 'Greedy function approximation: A gradient boosting machine,' *The Annals of Statistics*, vol. 29, no. 5, 2001-10. DOI: 10.1214/aos/1013203451.
 - [15] B. Bischl, *Introduction to machine learning (i2ml)*, Chapter 10.7: Bernoulli Loss, <https://introduction-to-machine-learning.netlify.app>, Accessed: 2022-01-29, 2022.
 - [16] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning*. Springer New York, 2009. DOI: 10.1007/978-0-387-84858-7.
 - [17] T. Parr and J. Howard, *How to explain gradient boosting*, Gradient boosting performs gradient descent, <https://explained.ai/gradient-boosting/descent.html>, Accessed: 2022-01-29, 2022.
 - [18] A. Subasi, *Practical machine learning for data analysis using python*. London: Academic Press, 2020, ISBN: 9780128213797.
 - [20] J. M. G. Christoph Schroer Felix Kruse, *A Systematic Literature Review on Applying CRISP-DM Process Model*. 2021.
 - [21] C. P. López, *DATA MINING. The CRISP-DM METHODOLOGY. The CLEM Language and IBM SPSS MODELER*. 2021-03-28.
 - [22] R. N. Gary D Miner Ken Yale, *Handbook of Statistical Analysis and Data Mining Applications*. 2017-11-09.
 - [23] M. Reddy, *API Design for C++*. Elsevier Science, 2011, ISBN: 9780123850041. [Online]. Available: <https://books.google.de/books?id=IY29LyIT85wC>.
 - [26] K. Lane, 'Intro to apis: What is an api?,' 2019-10-05. [Online]. Available: <https://blog.postman.com/intro-to-apis-what-is-an-api/> (visited on 2022-01-13).
 - [27] D. Gourley, B. Totty, M. Sayer, A. Aggarwal, and S. Reddy, *HTTP: The Definitive Guide* (Definitive Guides). O'Reilly Media, Incorporated, 2002, ISBN: 9781565925090.
 - [35] J. Hemming, *Methoden der Erforschung populärer Musik*. Gabler, Betriebswirt.-Vlg, 2015-10-28, 514 pp., ISBN: 3658114967. [Online]. Available: https://www.ebook.de/de/product/25204785/jan_hemming_methoden_der_erforschung_populaerer_musik.html.
 - [59] J. Davis and M. Goadrich, 'The relationship between precision-recall and ROC curves,' in *Proceedings of the 23rd international conference on Machine learning - ICML '06*, ACM Press, 2006. DOI: 10.1145/1143844.1143874.
 - [60] T. Fawcett, 'An introduction to roc analysis,' *Pattern recognition letters*, vol. 27, no. 8, pp. 861–874, 2006.
 - [61] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay, *Scikit-learn: Machine learning in Python*, Feature importances with a forest of trees, https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html, Accessed: 2022-01-29, 2011.

- [62] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay, *Scikit-learn: Machine learning in Python*, Feature importances with a forest of trees, https://scikit-learn.org/stable/auto_examples/ensemble/plot_forest_importances.html, Accessed: 2022-01-29, 2011.

Internet sources

- [19] 'Scikit learn user guide, 6.3. preprocessing data,' scikit-learn developers. (), [Online]. Available: <https://scikit-learn.org/stable/modules/preprocessing.html#standardization-or-mean-removal-and-variance-scaling>.
- [24] 'Introduction to web apis,' Mozilla. (), [Online]. Available: https://developer.mozilla.org/en-US/docs/Learn/JavaScript/Client-side_web_APIs/Introduction.
- [25] 'Types of apis & popular rest api protocol,' Stoplight. (), [Online]. Available: <https://stoplight.io/api-types/> (visited on 2022-01-24).
- [28] B. Cooksey. 'An introduction to apis.' B. Landers and D. Schreiber, Eds., Zapier, Inc. (2014-04-23), [Online]. Available: <https://zapier.com/learn/apis/chapter-1-introduction-to-apis/>.
- [29] 'Spotify web api documentation,' Spotify AB. (), [Online]. Available: <https://developer.spotify.com/documentation/web-api/>.
- [30] 'Working with json,' Mozilla. (), [Online]. Available: <https://developer.mozilla.org/en-US/docs/Learn/JavaScript/Objects/JSON>.
- [31] 'Json definiert,' Oracle. (), [Online]. Available: <https://www.oracle.com/de/database/what-is-json/>.
- [32] 'Introducing json.' (), [Online]. Available: <https://www.json.org/json-en.html> (visited on 2022-01-13).
- [33] R. Becker. 'What is music?' Study.com. (2021-10-29), [Online]. Available: <https://study.com/learn/lesson/what-is-music-characteristics-of-music-how-music-is-made.html>.
- [34] R. Havers. 'John coltrane quotes: The iconic saxophonist in his own words,' UDiscoverMusic. (2021-09-23), [Online]. Available: <https://www.udiscovermusic.com/stories/john-coltrane-in-20-quotes/>.
- [36] 'Chapter 2: Music: Fundamentals and educational roots in the u.s.' (), [Online]. Available: <https://milnepublishing.geneseo.edu/music-and-the-child/chapter/chapter-2/>.
- [37] 'Music theory: 5 fundamentals that you should know,' Shaw Academy. (2019-04), [Online]. Available: <https://www.shawacademy.com/blog/music-theory-5-fundamentals-that-you-should-know/>.
- [38] M. Shipman. 'Analysis of 50 years of hit songs yields tips for advertisers,' NC State University. (2014-03-18), [Online]. Available: <https://news.ncsu.edu/2014/03/wms-henard-hits2014/>.
- [39] 'What are we listening to? | the most popular themes in songs,' Atlanta Institute of Music and Media. (2020-02-14), [Online]. Available: <https://www.aimm.edu/blog/most-popular-themes-in-songs>.
- [40] 'Instrument classification,' Goshen College. (), [Online]. Available: <https://www.goshen.edu/academics/music/mary-k-oyer-african-music-archive/instrument-classification/>.

-
- [41] 'Musikrichtungen,' MUSICFLX. (), [Online]. Available: <https://www.musicflx.de/musikrichtungen/>.
 - [42] 'What is hip-hop?' Musical Dictionary. (), [Online]. Available: <https://musicaldictionary.com/what-is-hip-hop/>.
 - [43] G. Tate. 'Hip-hop music and cultural movement,' Encyclopædia Britannica, Inc. (), [Online]. Available: <https://www.britannica.com/art/hip-hop>.
 - [44] R. PQ. 'Hip hop history: From the streets to the mainstream,' Icon Collective. (2019-11-13), [Online]. Available: <https://iconcollective.edu/hip-hop-history/>.
 - [45] K. Goodrich. 'What is hip-hop?' (2017-04-25), [Online]. Available: <https://medium.com/@katiegoodrich/what-is-hip-hop-9f9abfffb25e>.
 - [46] M. Beek. 'What is jazz music?' BBC Music Magazine. (2021-07-15), [Online]. Available: <https://www.classical-music.com/features/articles/jazz-music-what-it-is-and-how-it-evolved/>.
 - [47] D. J. Wildridge. 'Characteristics of jazz: An introduction,' CMUSE. (2020-01-15), [Online]. Available: [BBC%20Music%20Magazine](https://www.bbc.com/music/articles/characteristics-of-jazz-an-introduction-20200115).
 - [48] 'What is jazz?' National Museum of American History. (), [Online]. Available: <https://americanhistory.si.edu/smithsonian-jazz/education/what-jazz>.
 - [49] 'What is jazz? a guide to the history and sound of jazz,' MasterClass. (2020-11-08), [Online]. Available: <https://www.masterclass.com/articles/what-is-jazz#what-is-jazz-music>.
 - [50] 'What is jazz?' Musical Dictionary. (), [Online]. Available: <https://musicaldictionary.com/what-is-jazz/>.
 - [51] 'Rock 'n' roll music guide: 4 characteristics of rock 'n' roll,' MasterClass. (2021-10-25), [Online]. Available: <https://www.masterclass.com/articles/rock-n-roll-music-guide>.
 - [52] 'What is rock music?' Musical Dictionary. (), [Online]. Available: <https://musicaldictionary.com/what-is-rock-music/>.
 - [53] B. Clark. '30 different types of rock music (rock subgenres),' MusicianWave.com. (2021-08-19), [Online]. Available: <https://www.musicianwave.com/types-of-rock-music-subgenres/>.
 - [54] 'Spotify authorization documentation,' Spotify AB. (), [Online]. Available: <https://developer.spotify.com/>.
 - [55] L.Rabe. 'Spotify: Daten und fakten zum schwedischen musikstreaming-anbieter.' (2021-10-14), [Online]. Available: <https://de.statista.com/themen/1894/spotify/#dossierKeyfigures>.
 - [56] B. Stephenson. 'What is spotify and how does it work?' (2021-04-28), [Online]. Available: <https://www.lifewire.com/what-is-spotify-4685829>.
 - [57] C. Boyd. 'How spotify recommends your new favorite artist.' (2019-11-11), [Online]. Available: <https://towardsdatascience.com/how-spotify-recommends-your-new-favorite-artist-8c1850512af0>.

- [58] G. T. A. Features, Ed. (), [Online]. Available: <https://developer.spotify.com/documentation/web-api/reference/#/operations/get-audio-features>.

Declaration in lieu of oath

I hereby declare that I produced the submitted paper with no assistance from any other party and without the use of any unauthorized aids and, in particular, that I have marked as quotations all passages which are reproduced verbatim or near-verbatim from publications. Also, I declare that the submitted print version of this thesis is identical with its digital version. Further, I declare that this thesis has never been submitted before to any examination board in either its present form or in any other similar version. I herewith **agree/disagree** that this thesis may be published. I herewith consent that this thesis may be uploaded to the server of external contractors for the purpose of submitting it to the contractors' plagiarism detection systems. Uploading this thesis for the purpose of submitting it to plagiarism detection systems is not a form of publication.

Düsseldorf, 31.1.2022

(Location, Date)

A handwritten signature in black ink, consisting of a large, stylized 'H' followed by a series of loops and a final flourish.

(handwritten signature)